

# Kubernetes Security Concepts

Daniel Hinojosa

Fast Guide to Securing your Cluster



# In this Presentation

- Review of Kubernetes Architecture
- Transport Layer Security (TLS)
  - mTLS within Kubernetes
  - Setting up TLS Termination
- Users and Roles
- Role Bindings and Cluster Bindings
- Authentication
- Kubernetes Network Policies

**Slides and Material:** <https://github.com/dhinojosa/k8s-security-concepts>

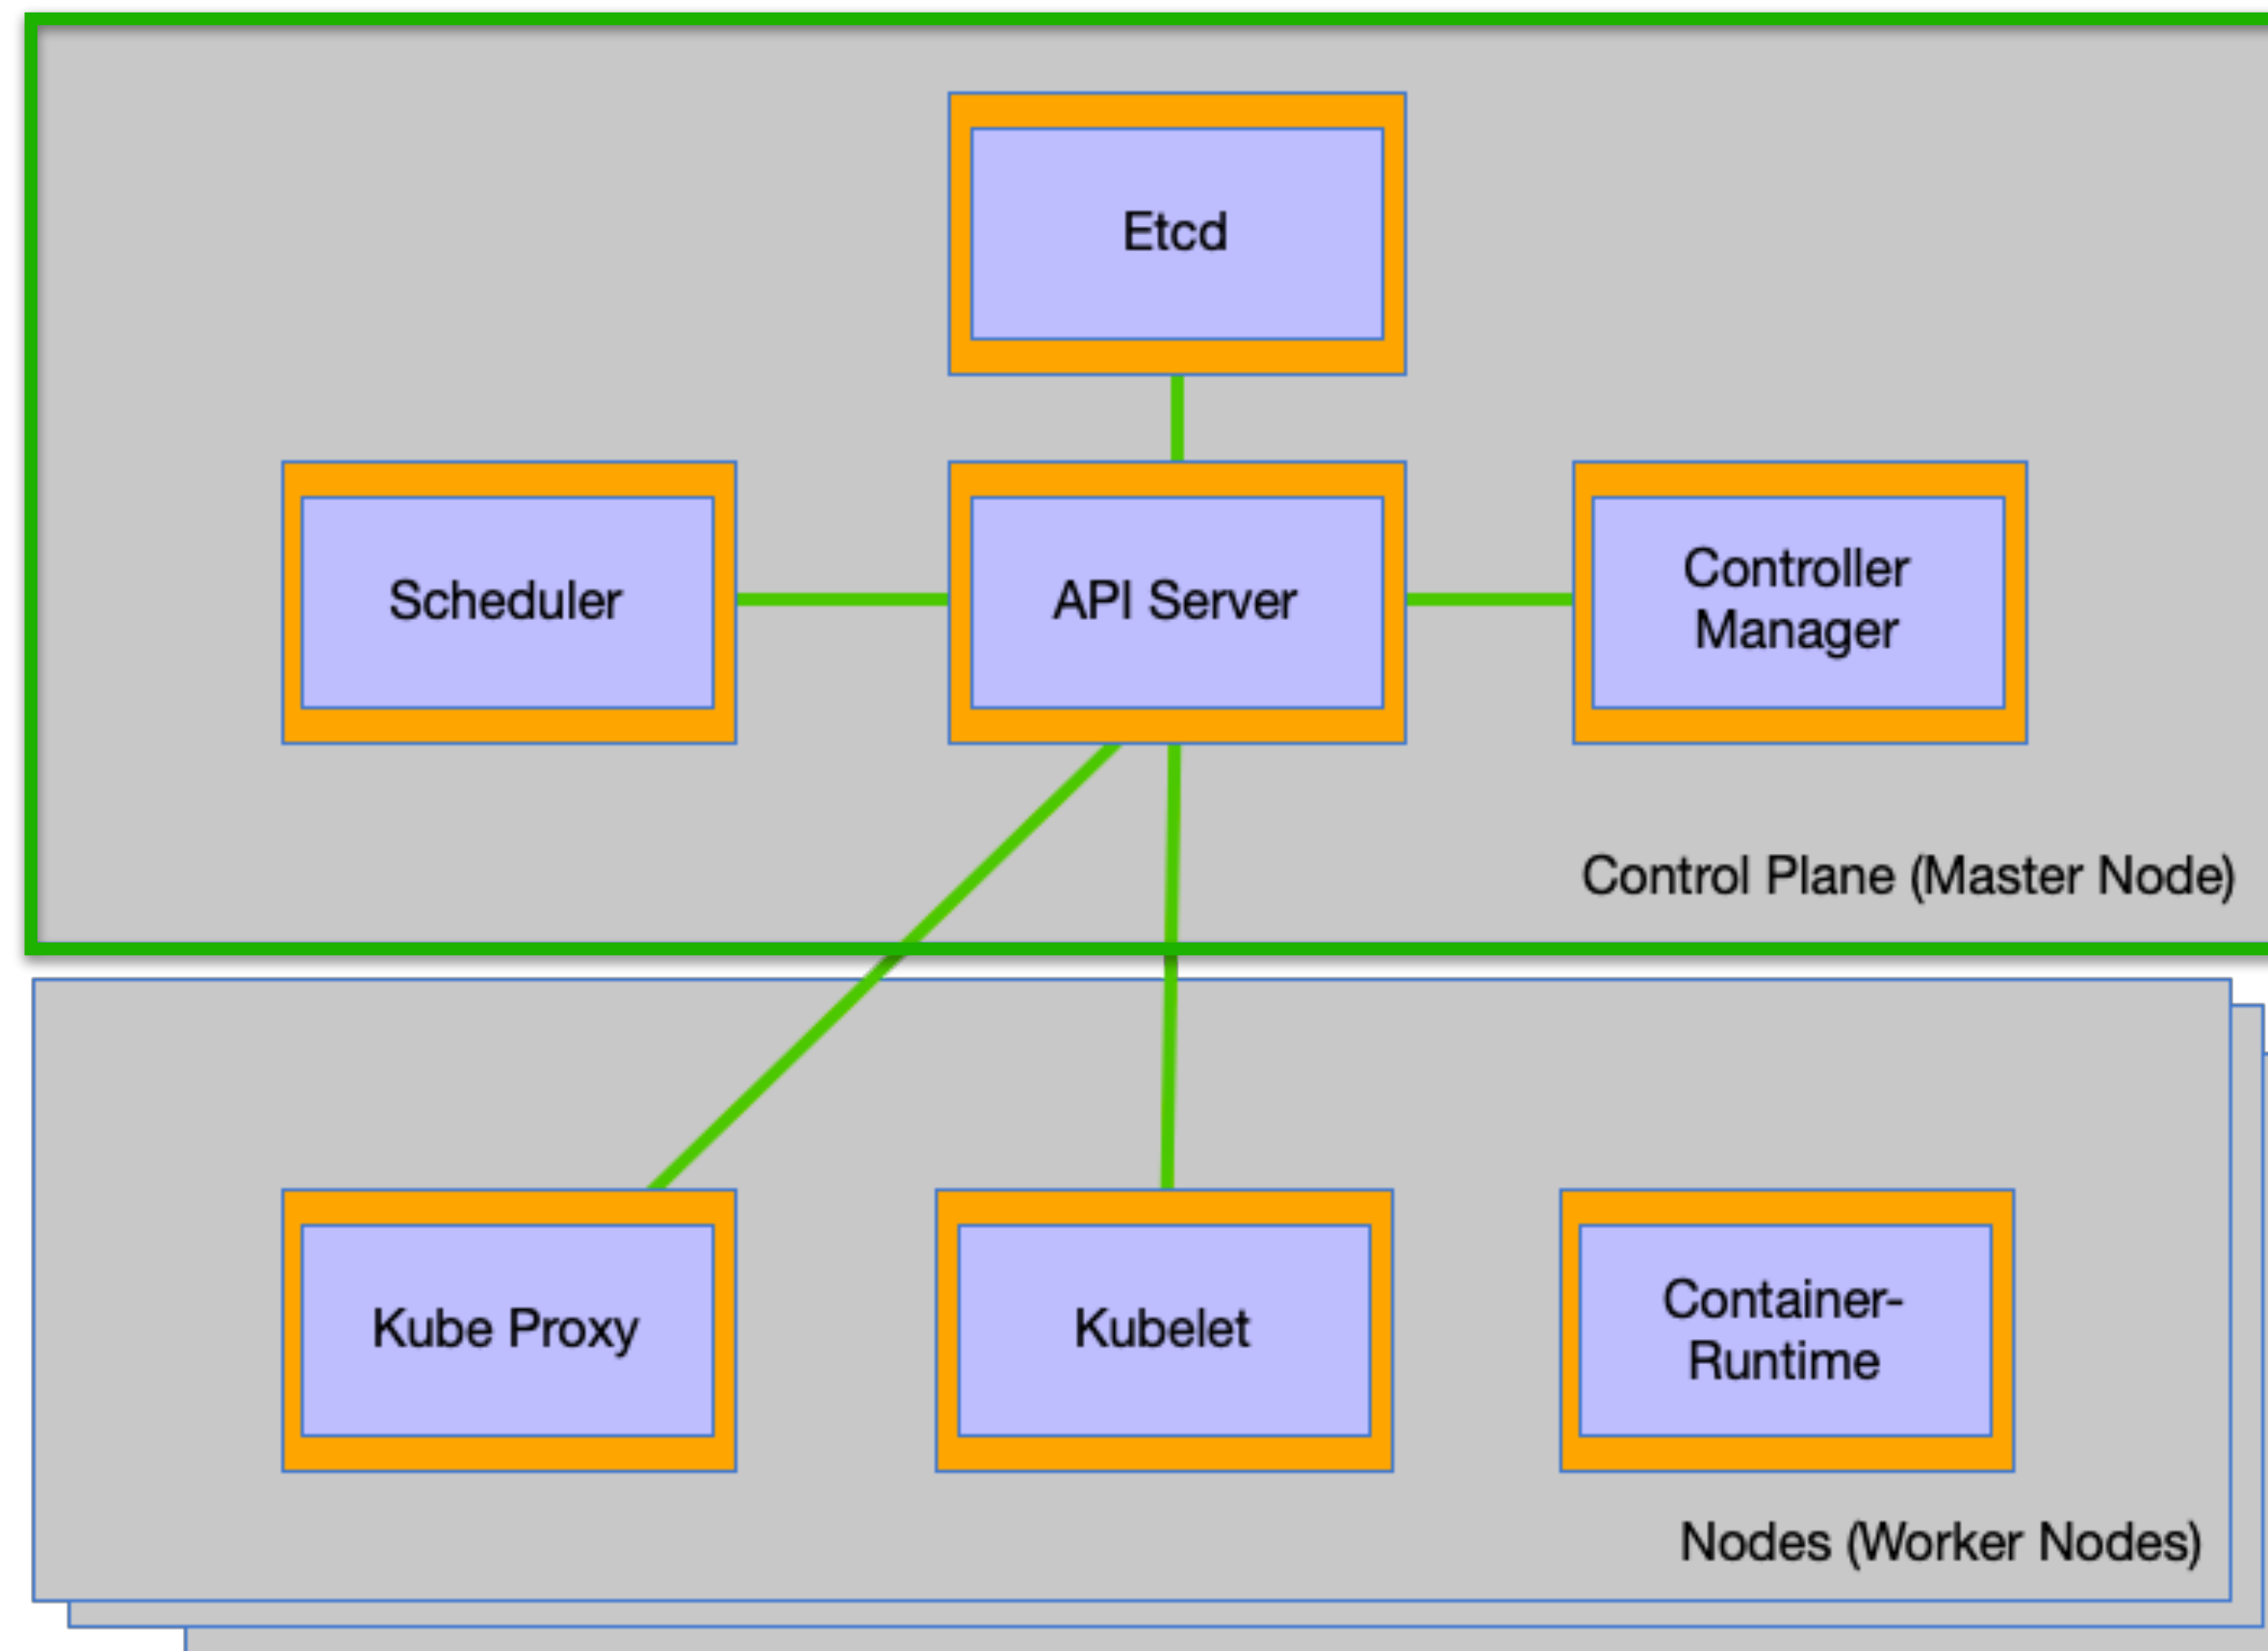


# Architecture Overview



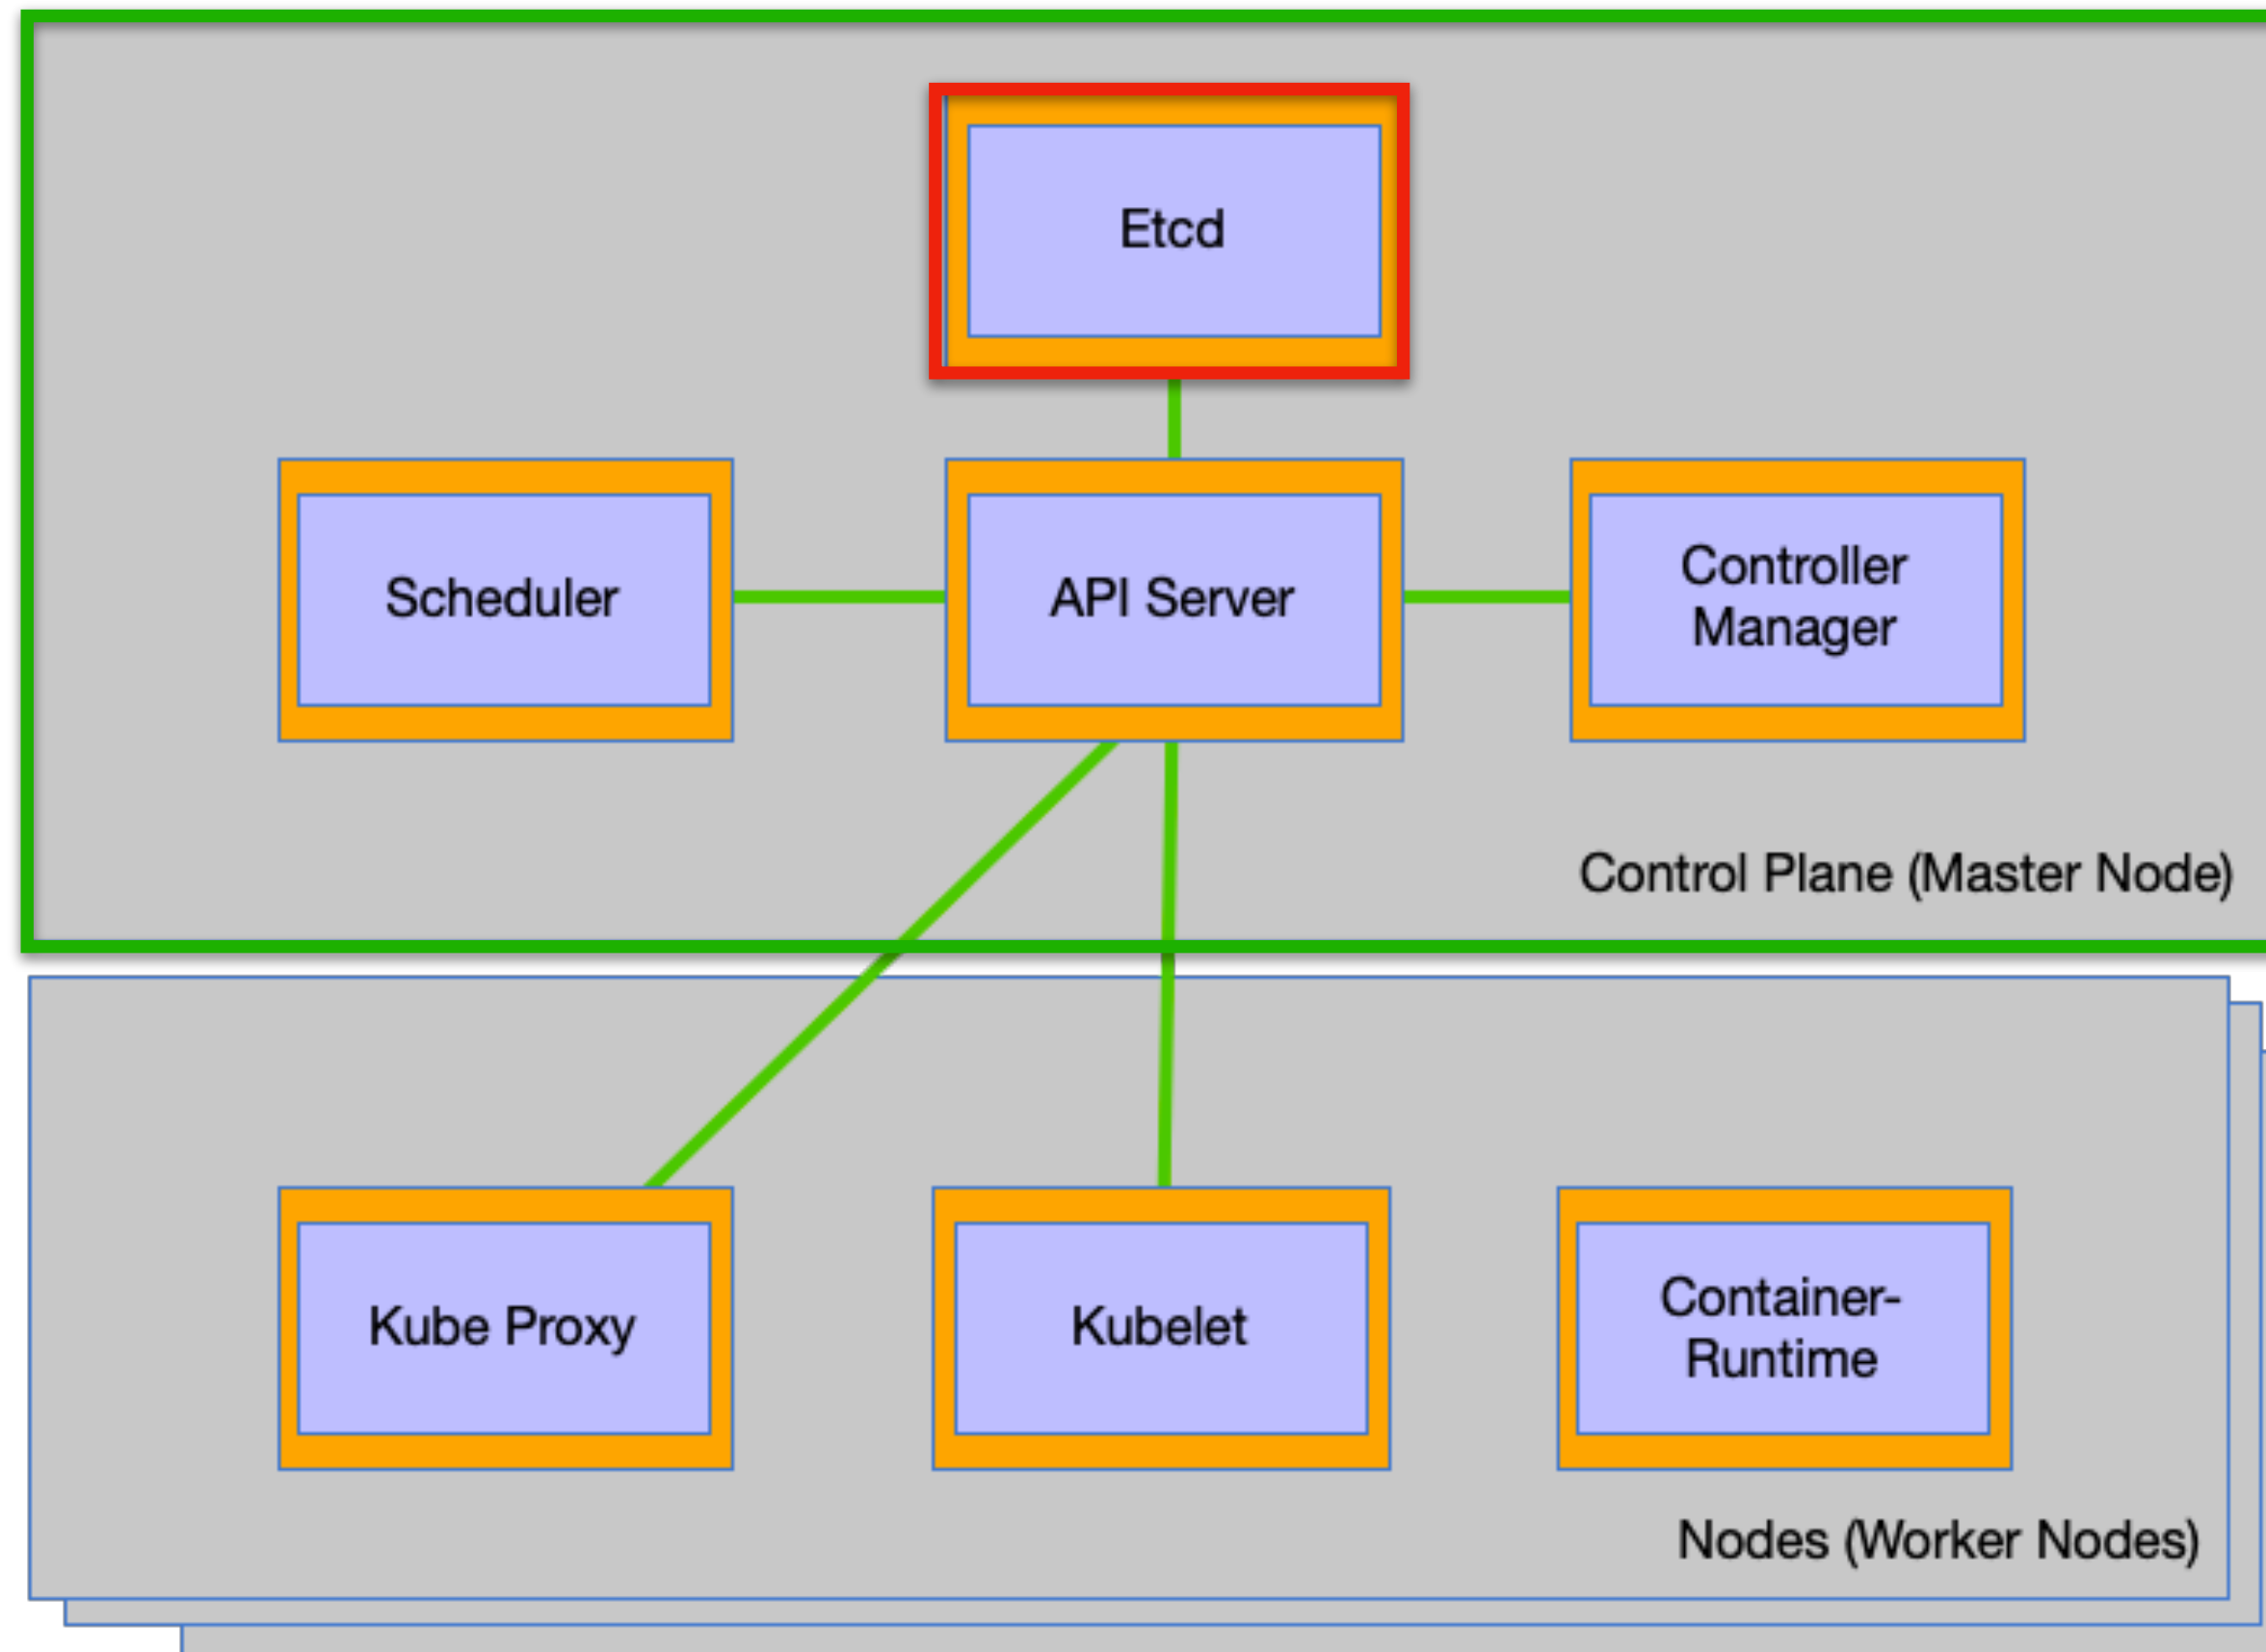


# Control Plane



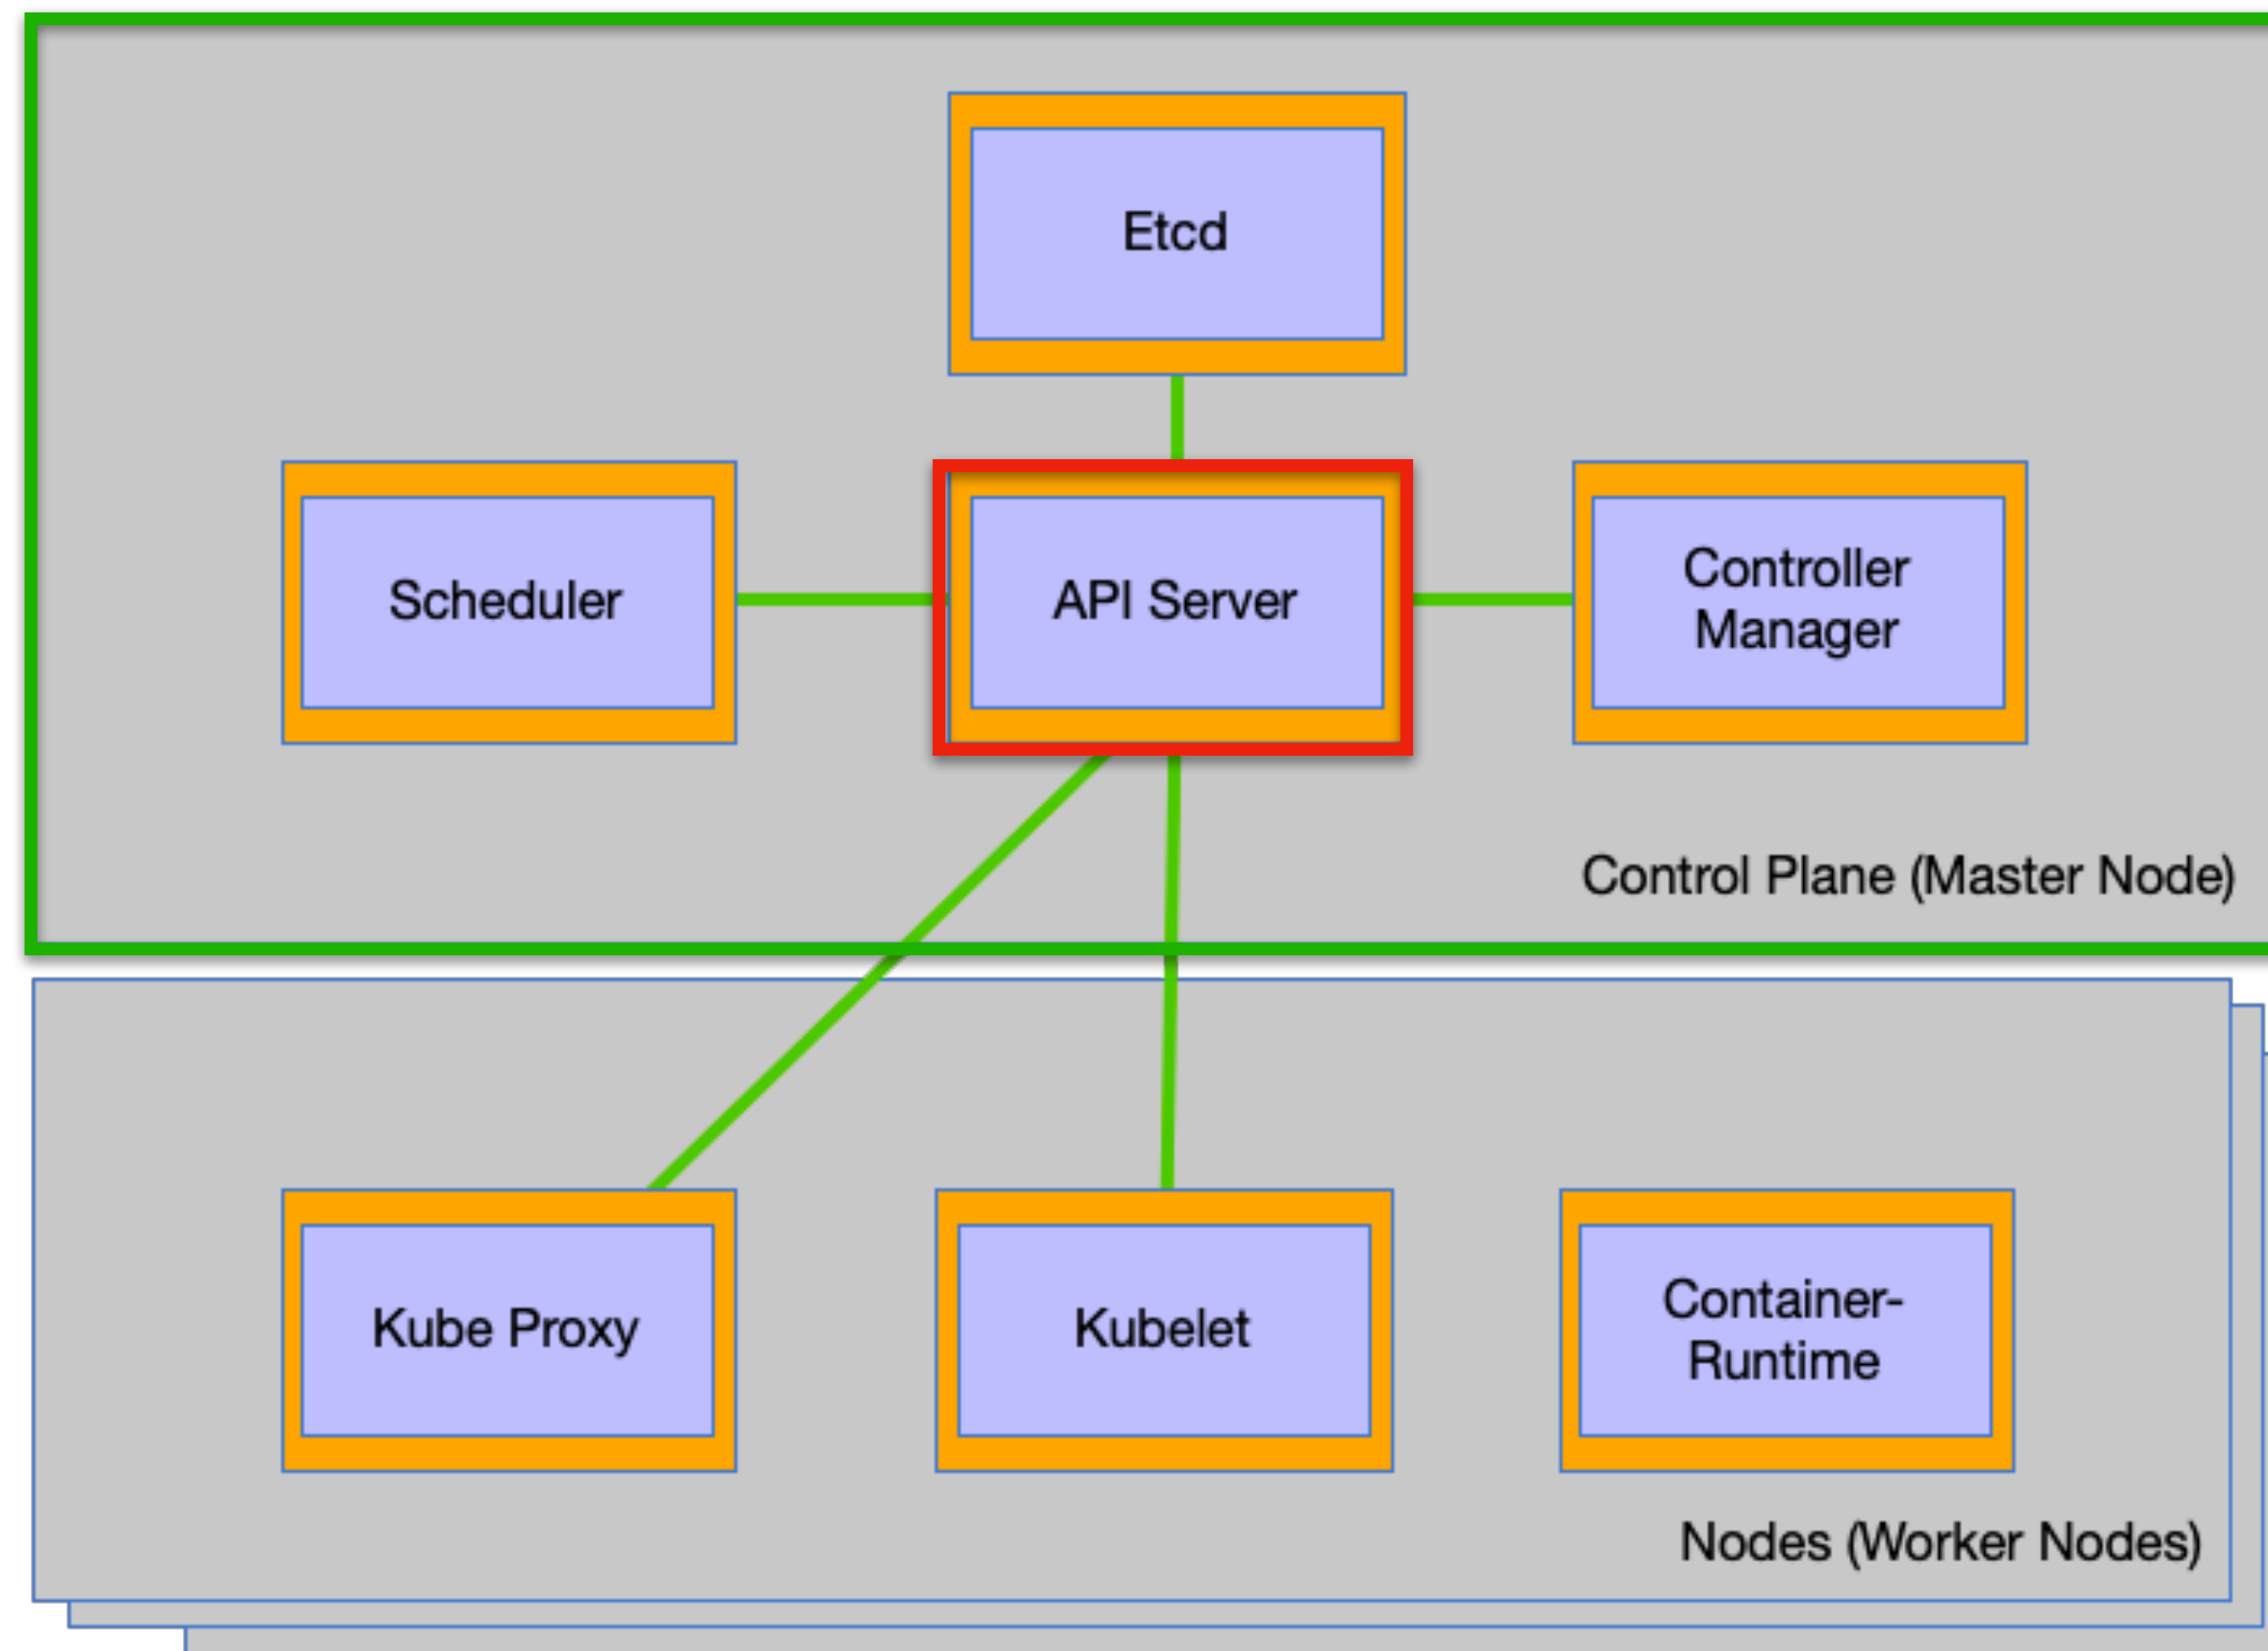
- Makes the whole cluster function
- Components included
  - Etcd Storage
  - API Server
  - Scheduler
  - Controller Manager

# Etc



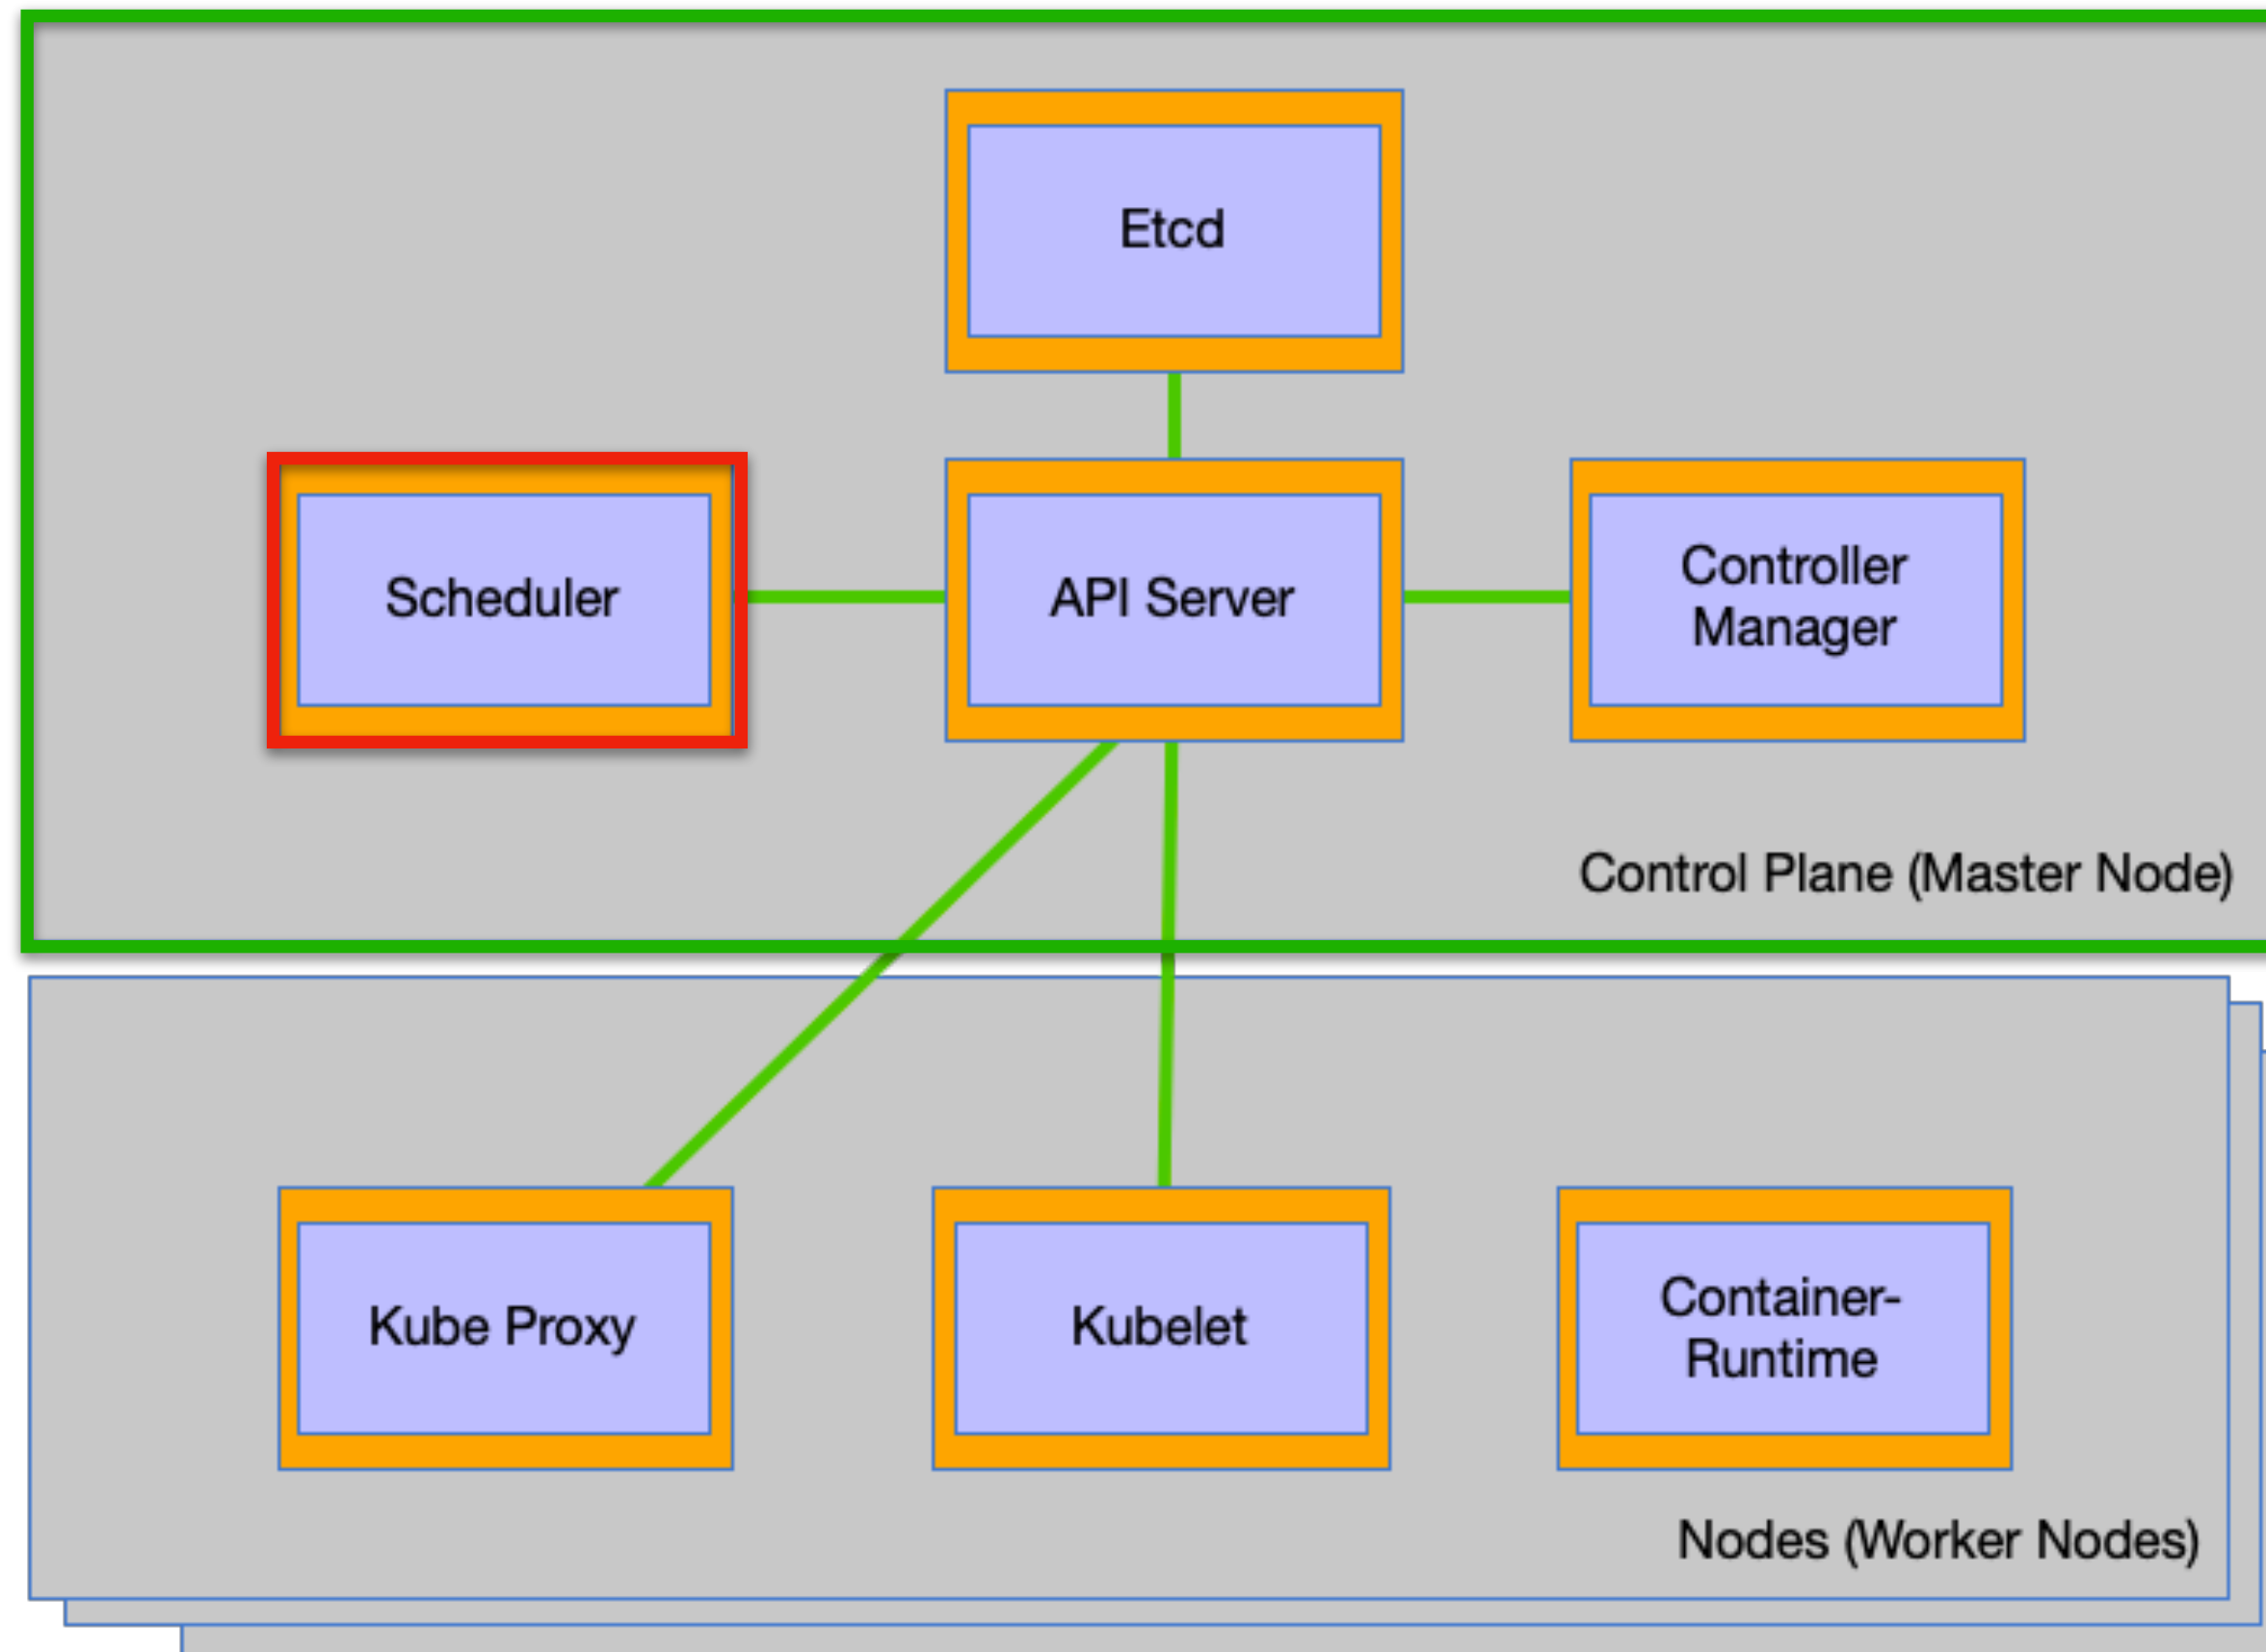
- Fast, Distributed, Key-Value Store
- Component Manifests are stored in etc
- More than one instance can be used for High Availability
- All read-writes are done through the API Server
- Only component that stores state and metadata

# API Server



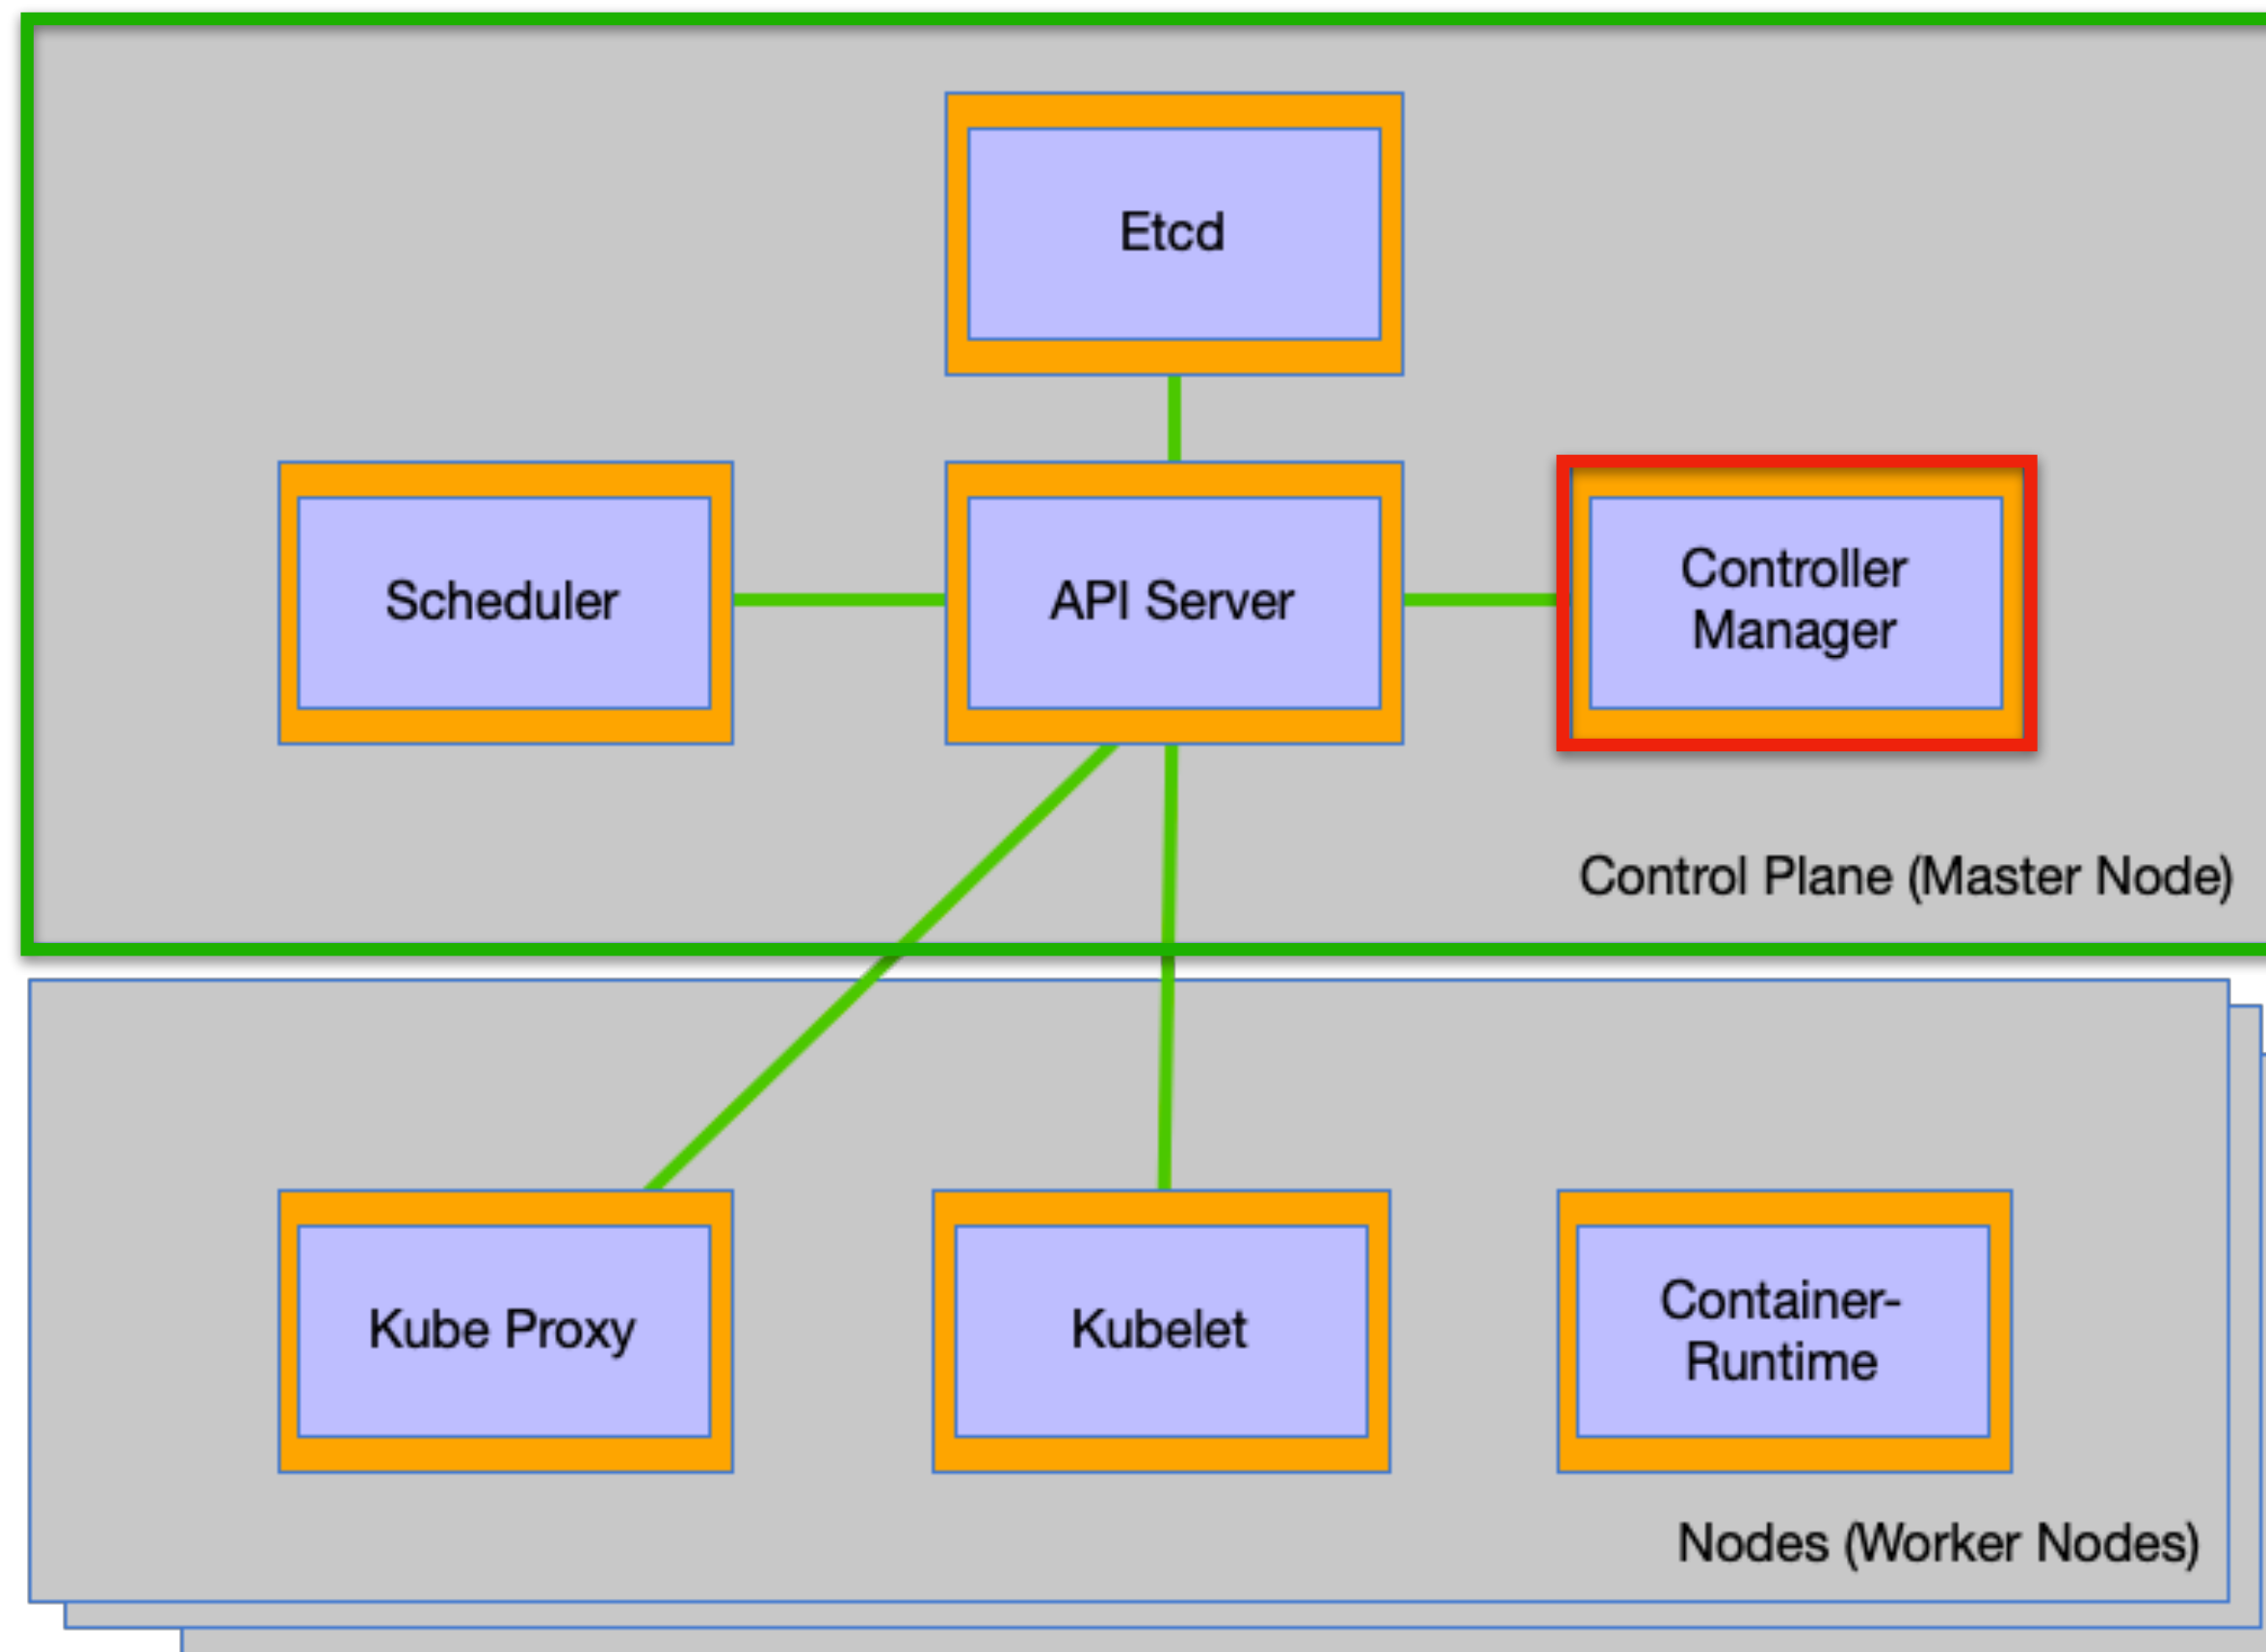
- RestFUL server used by other components
- All state to the API Server
- Embedded Validation: Components cannot store invalid key-value data
- Optimistic Locking: Changes to an object are not overridden
- Notify Clients of their change

# Scheduler



- Wait for Pod Events from the API Server as a Watch
- Assign a Node to Each Pod
- But it does not run the Pod
- It merely schedules the pod to run on its node
- It does so by sending the pod schedule to the the API Server

# Controller Manager

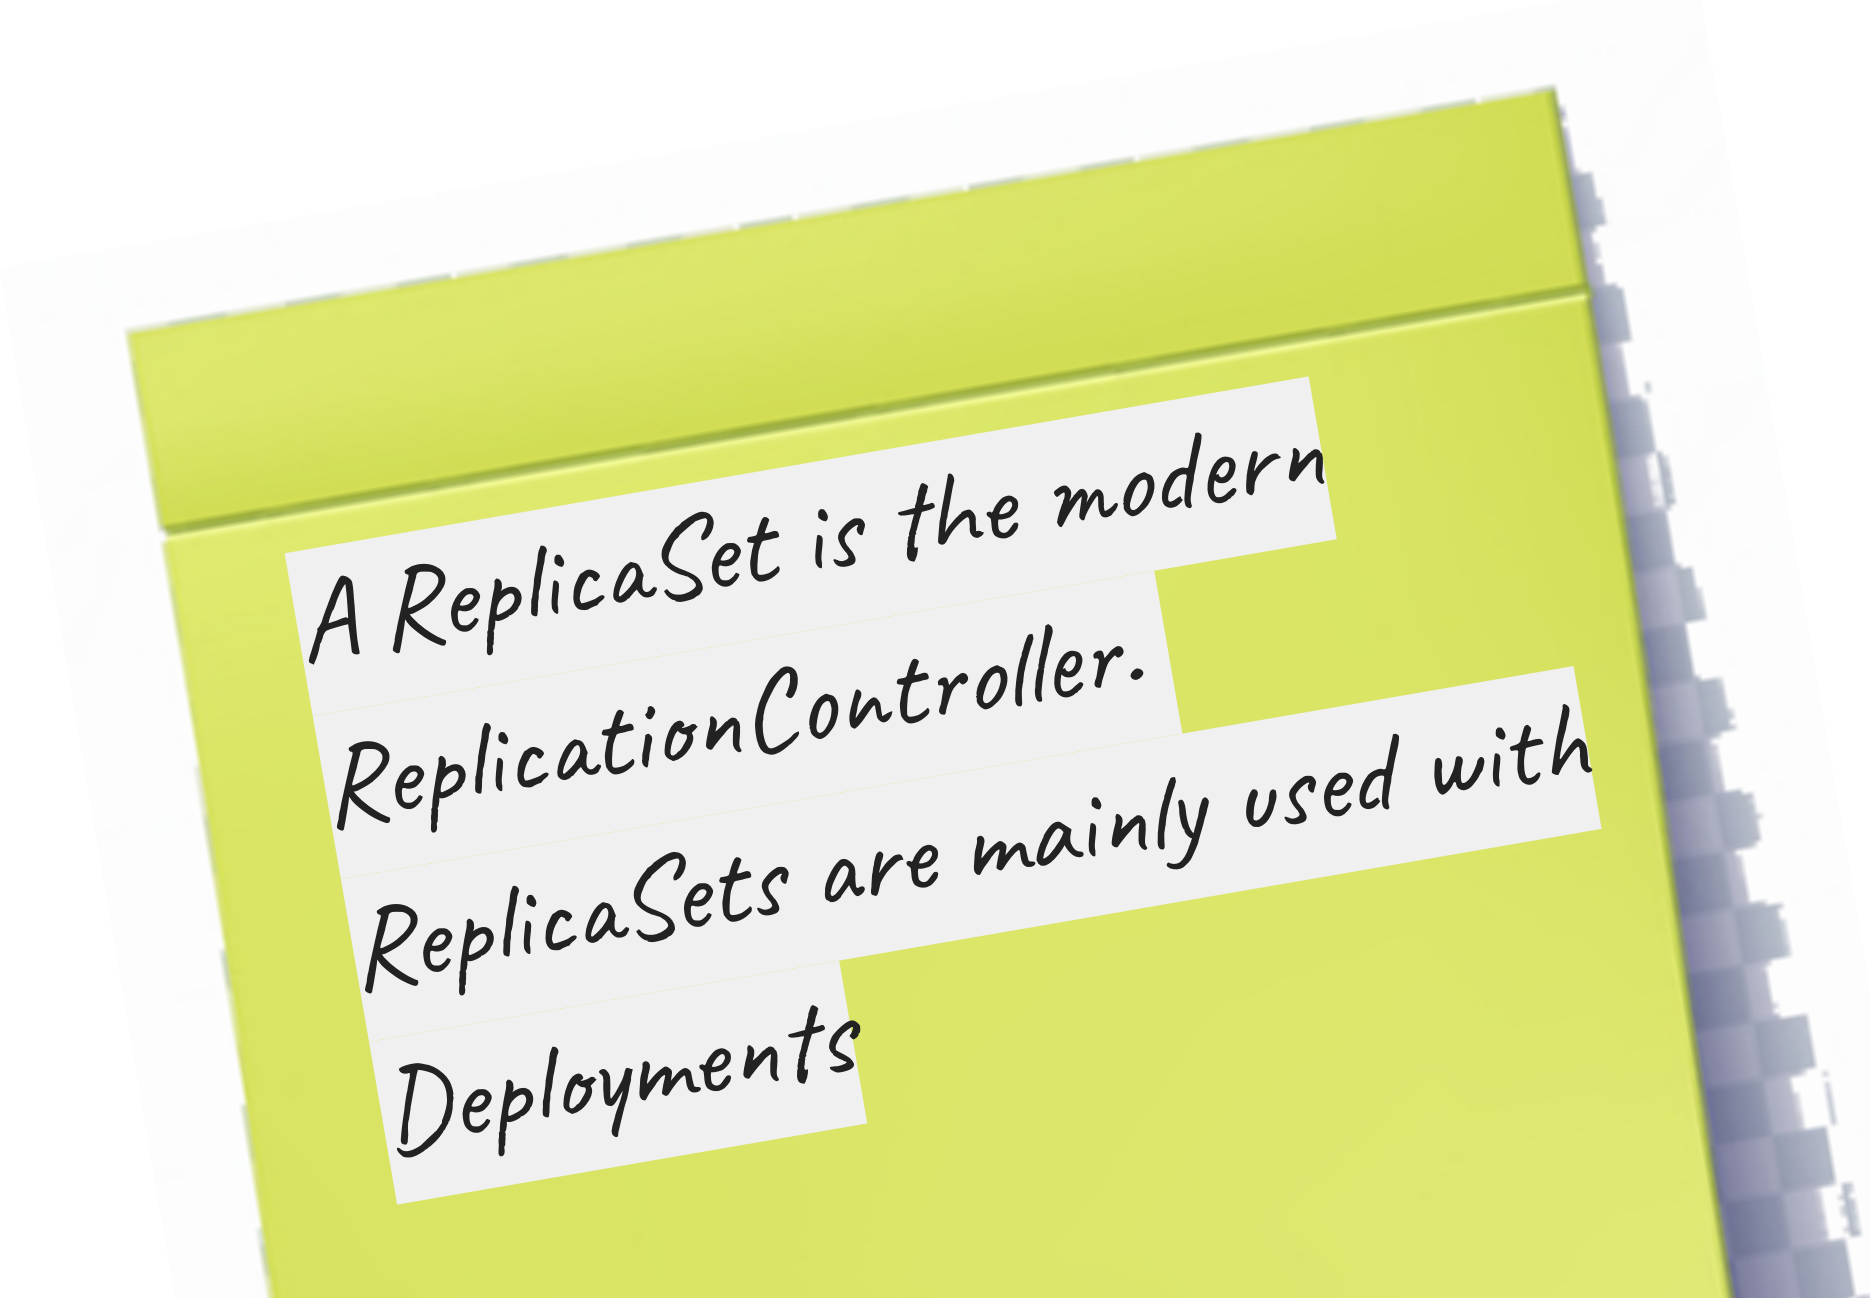


- Component that ensures that the state of the system converges towards the correct state
- Combines multiple *controllers* to perform reconciliation
- Controllers are spawned in multiple processes
- Has the ability to replace with a custom controller



# List of Varying Controllers

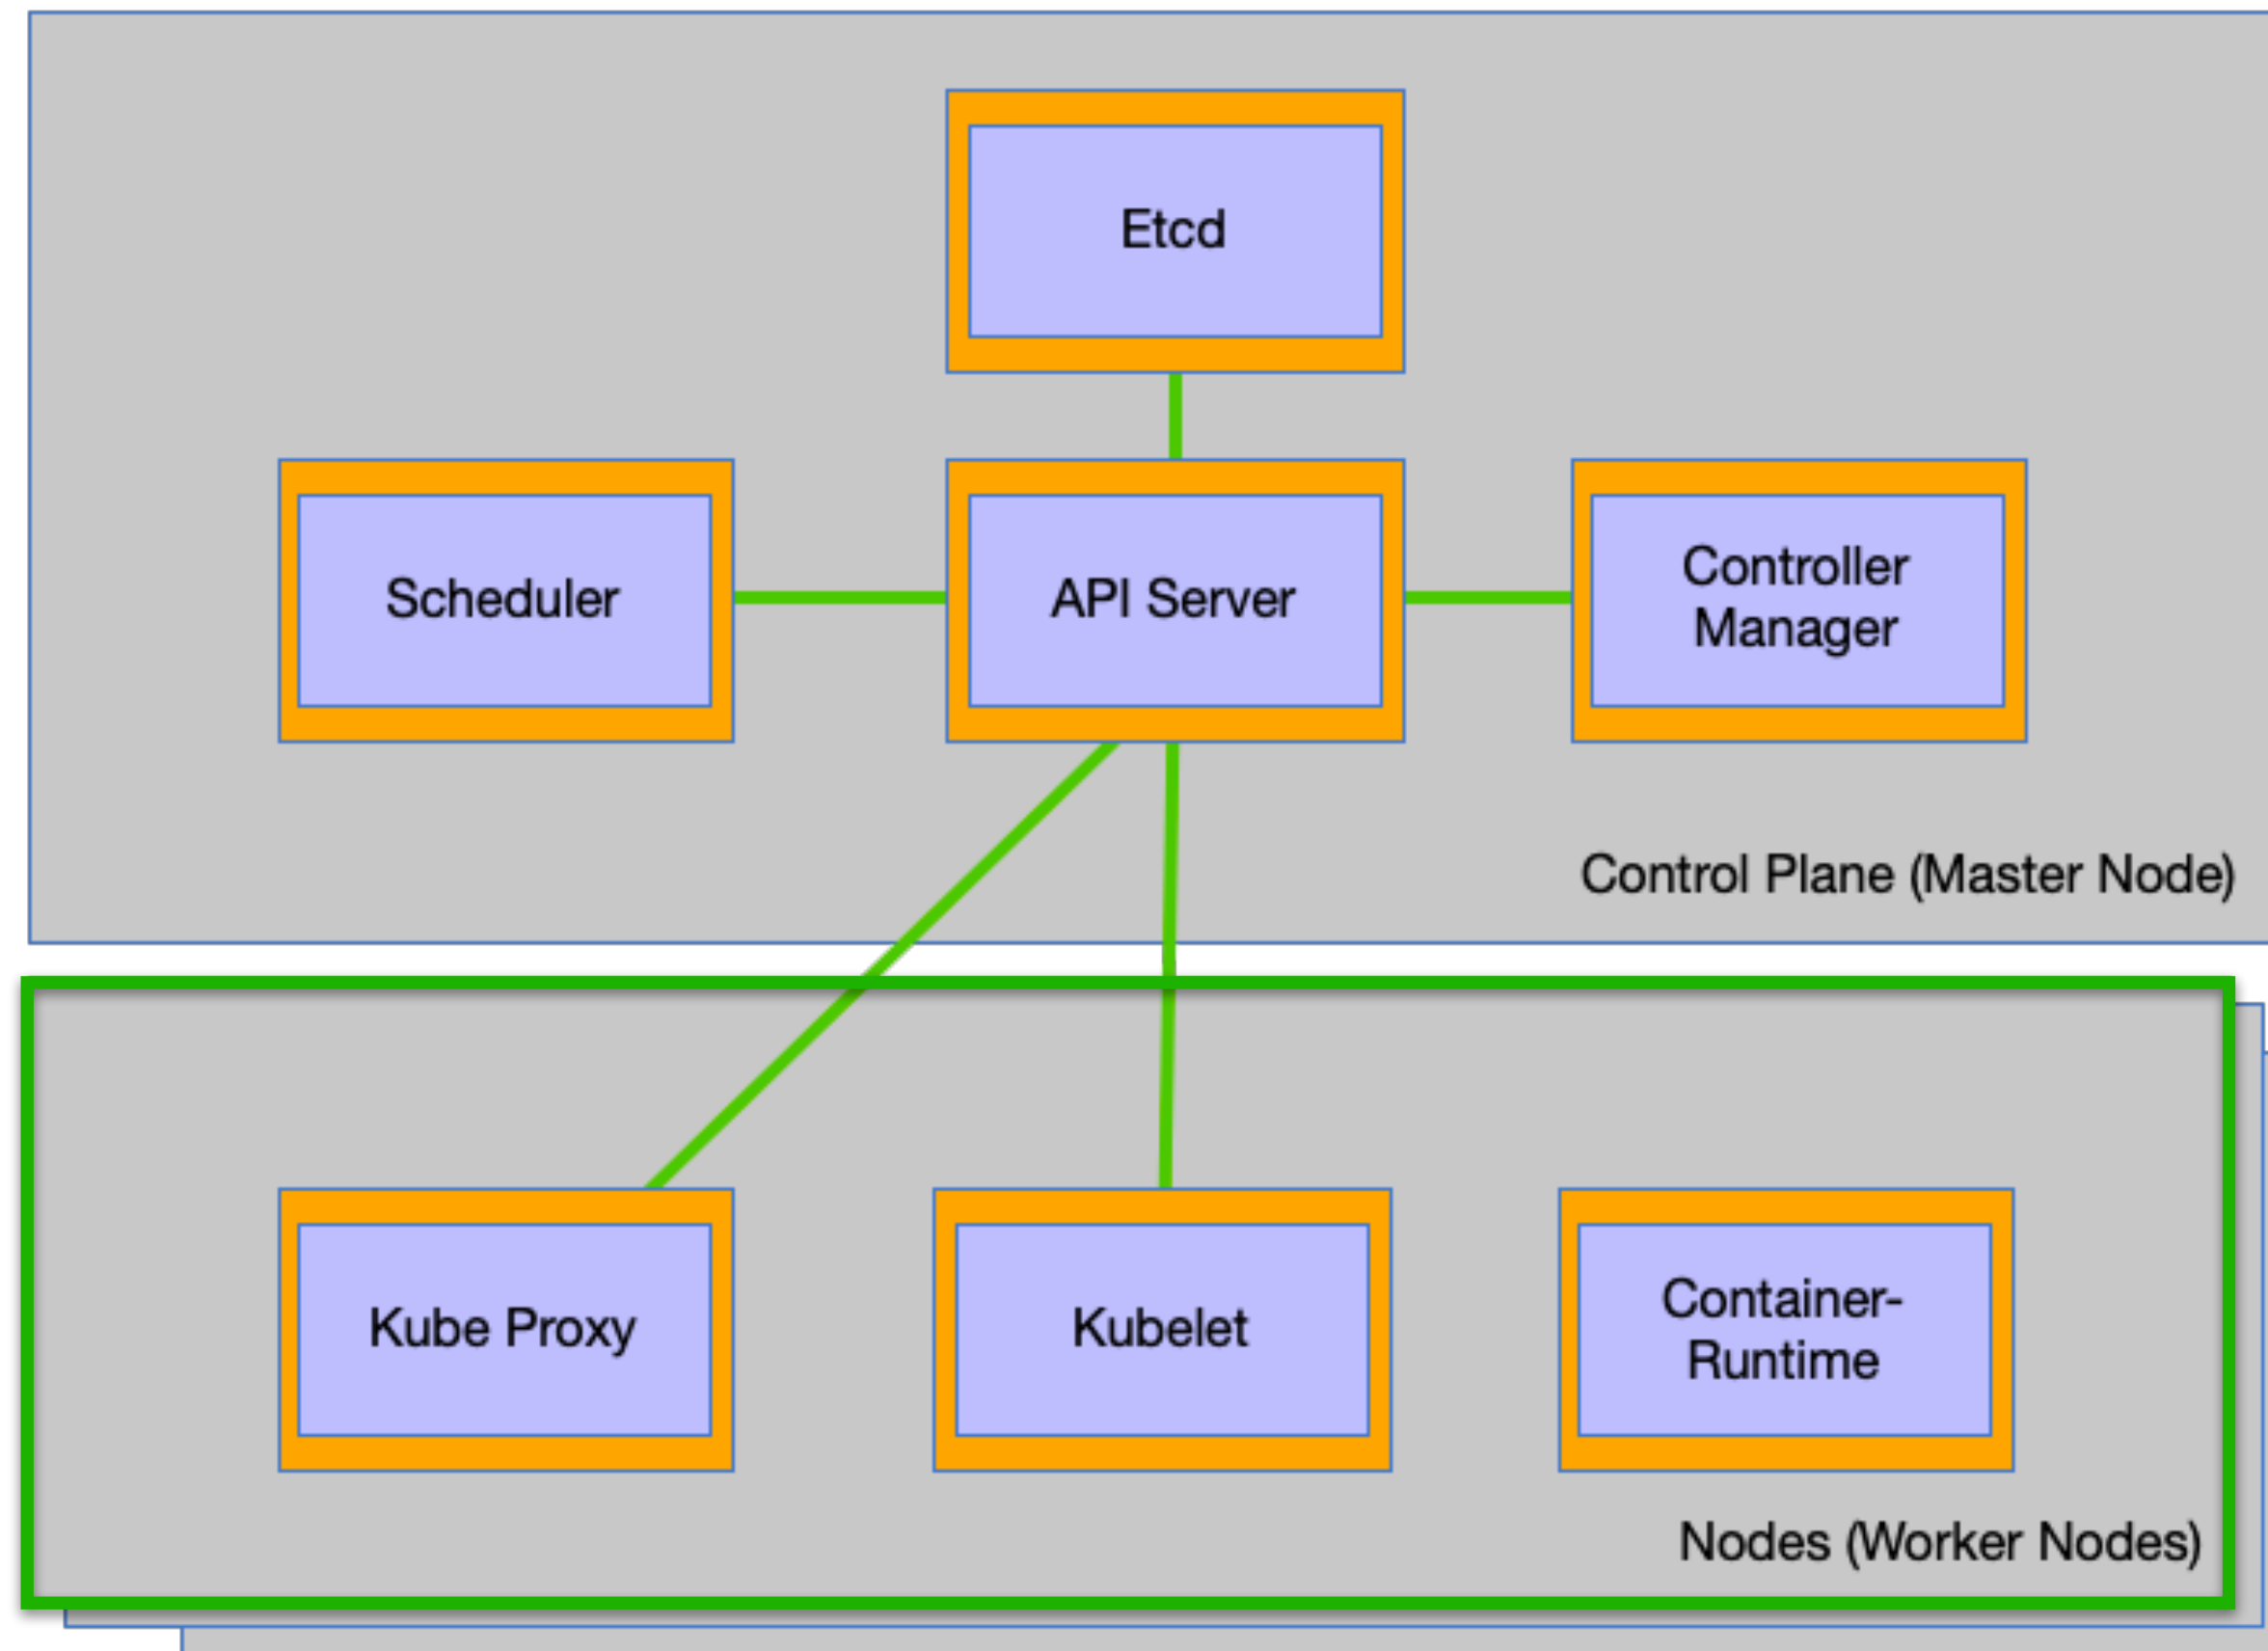
- Service Controller
- Persistent Volume Controller
- Replica Set Controller
- Daemon Set Controller
- Job Controller
- Stateful Set Controller
- Node Controller
- Service Controller
- Endpoints Controller
- Namespace Controller
- Replication Controller Controller



A ReplicaSet is the modern  
ReplicationController.  
ReplicaSets are mainly used with  
Deployments



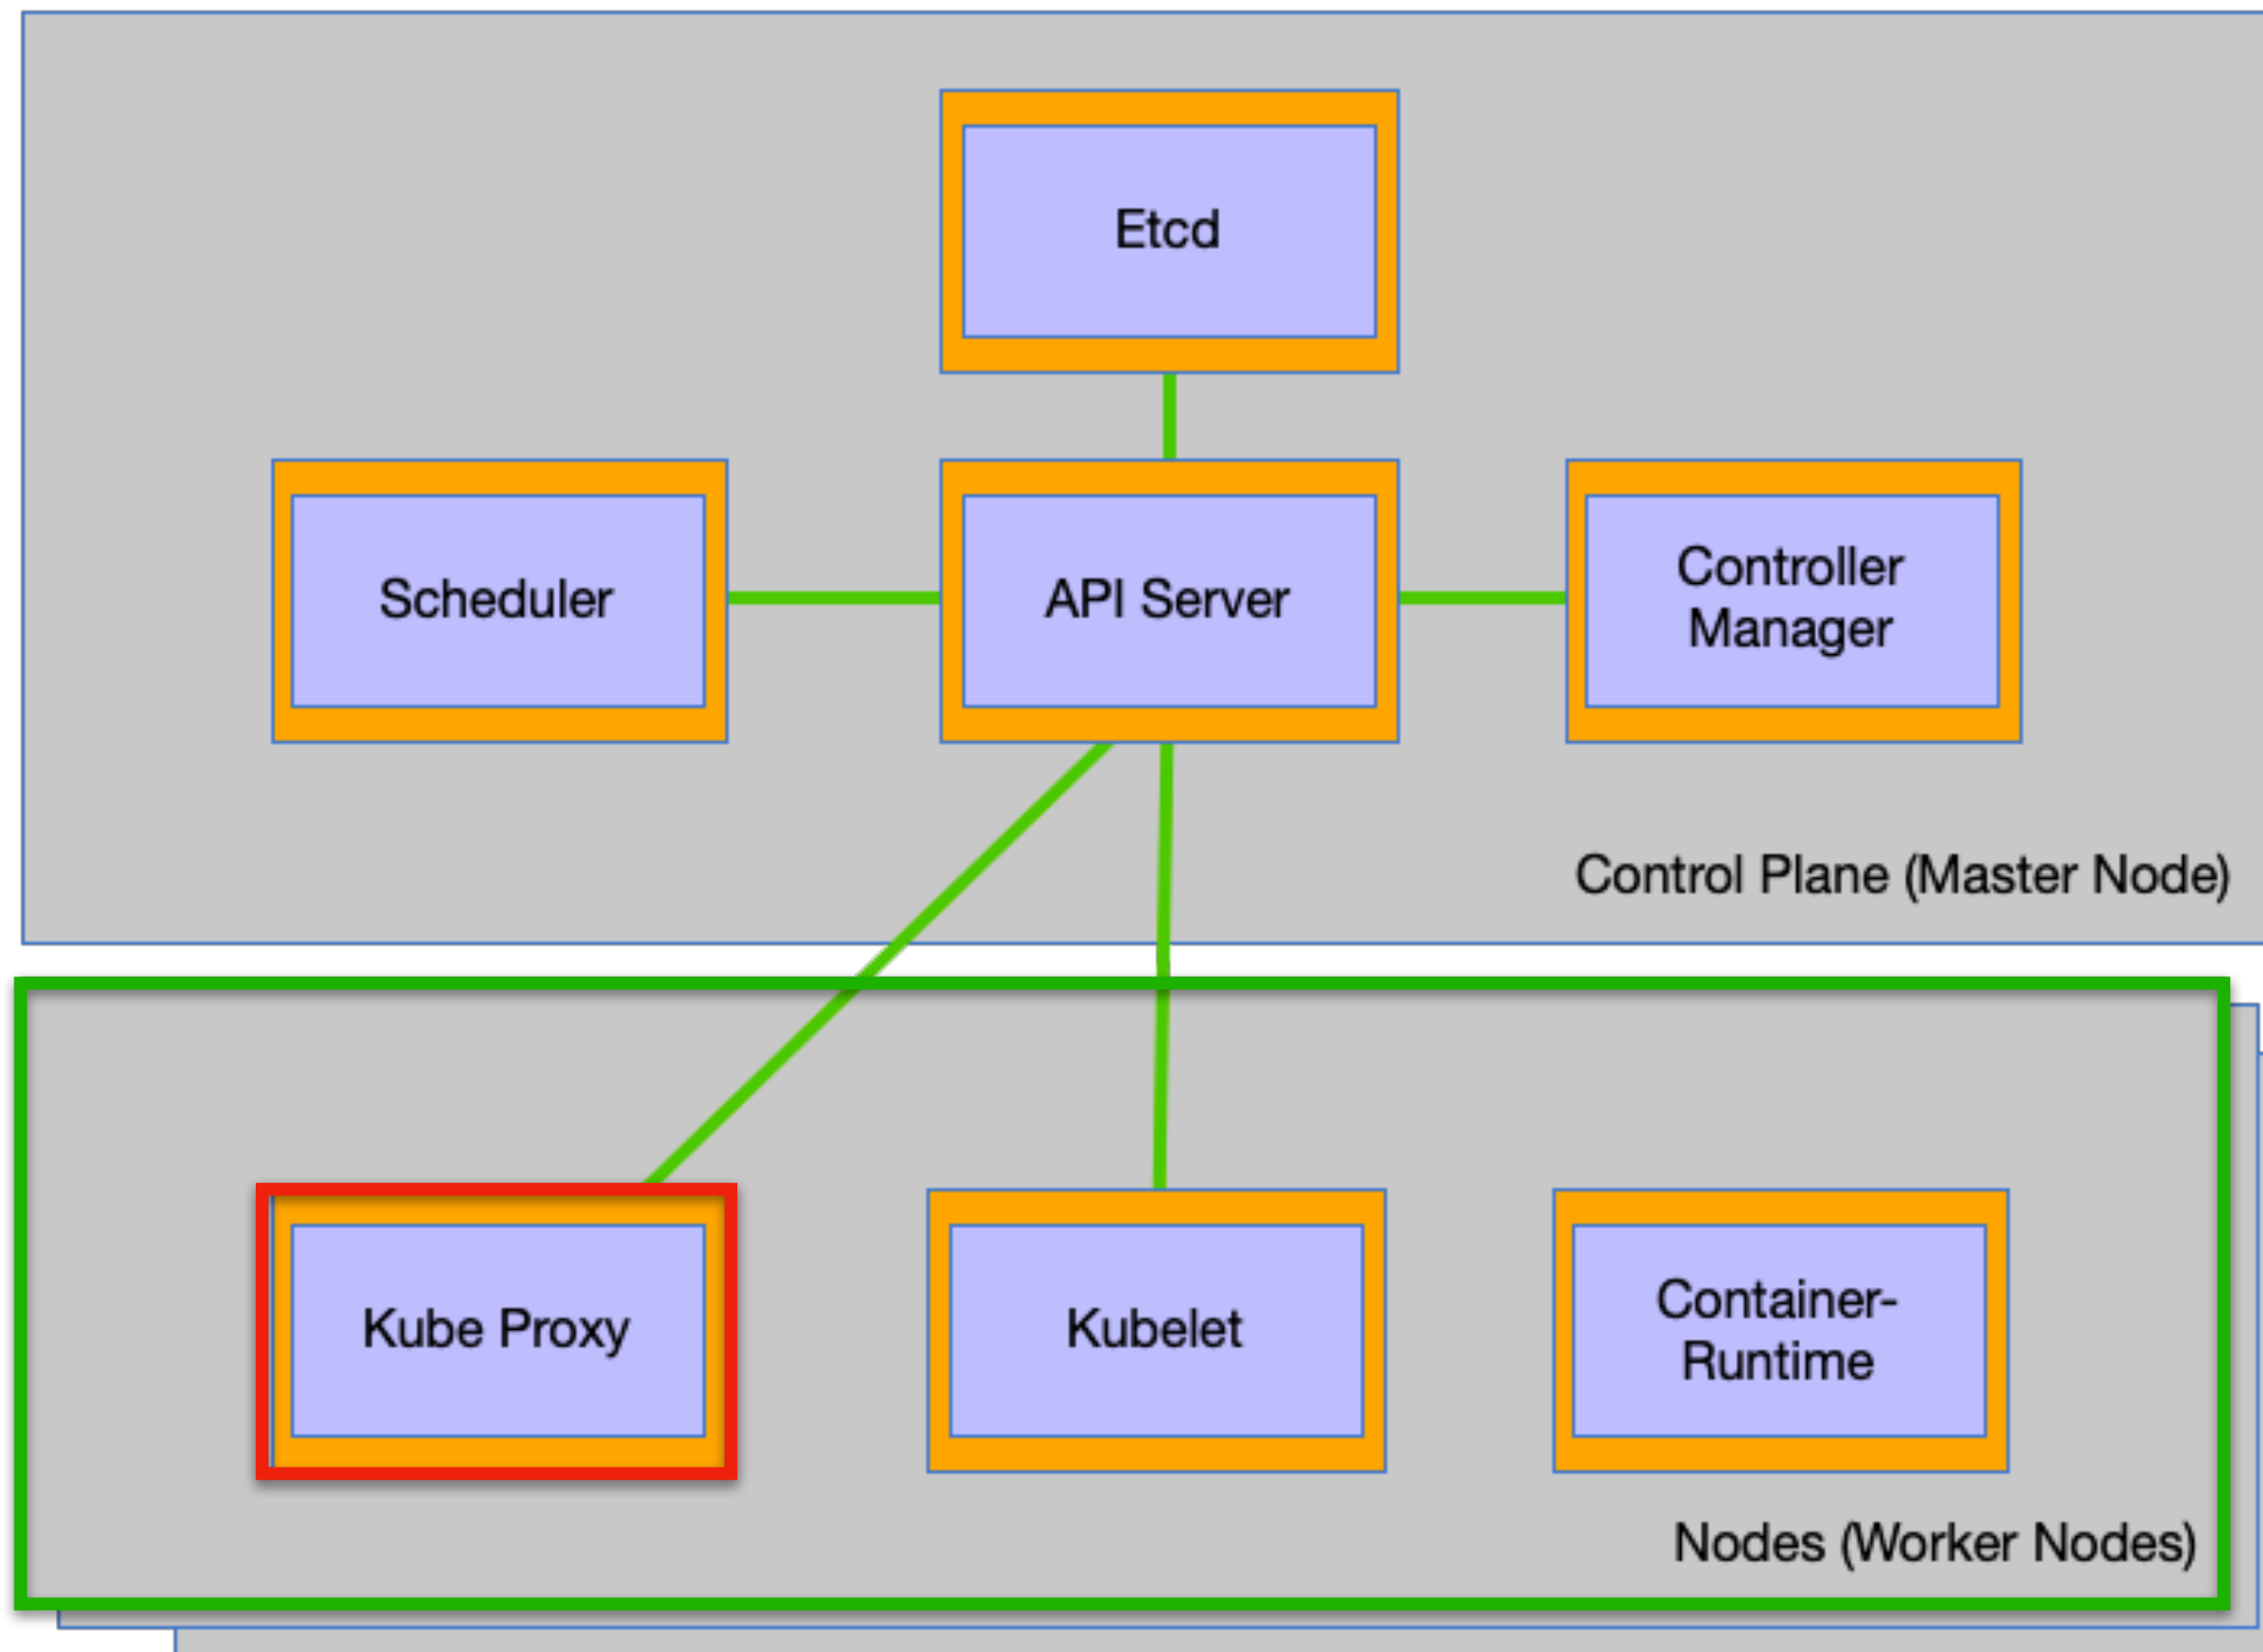
# Node



- A node may be a virtual or physical machine
- Each node contains the services necessary to run Pods
- Contains
  - Kube Proxy
  - Kubelet
  - Container Runtime



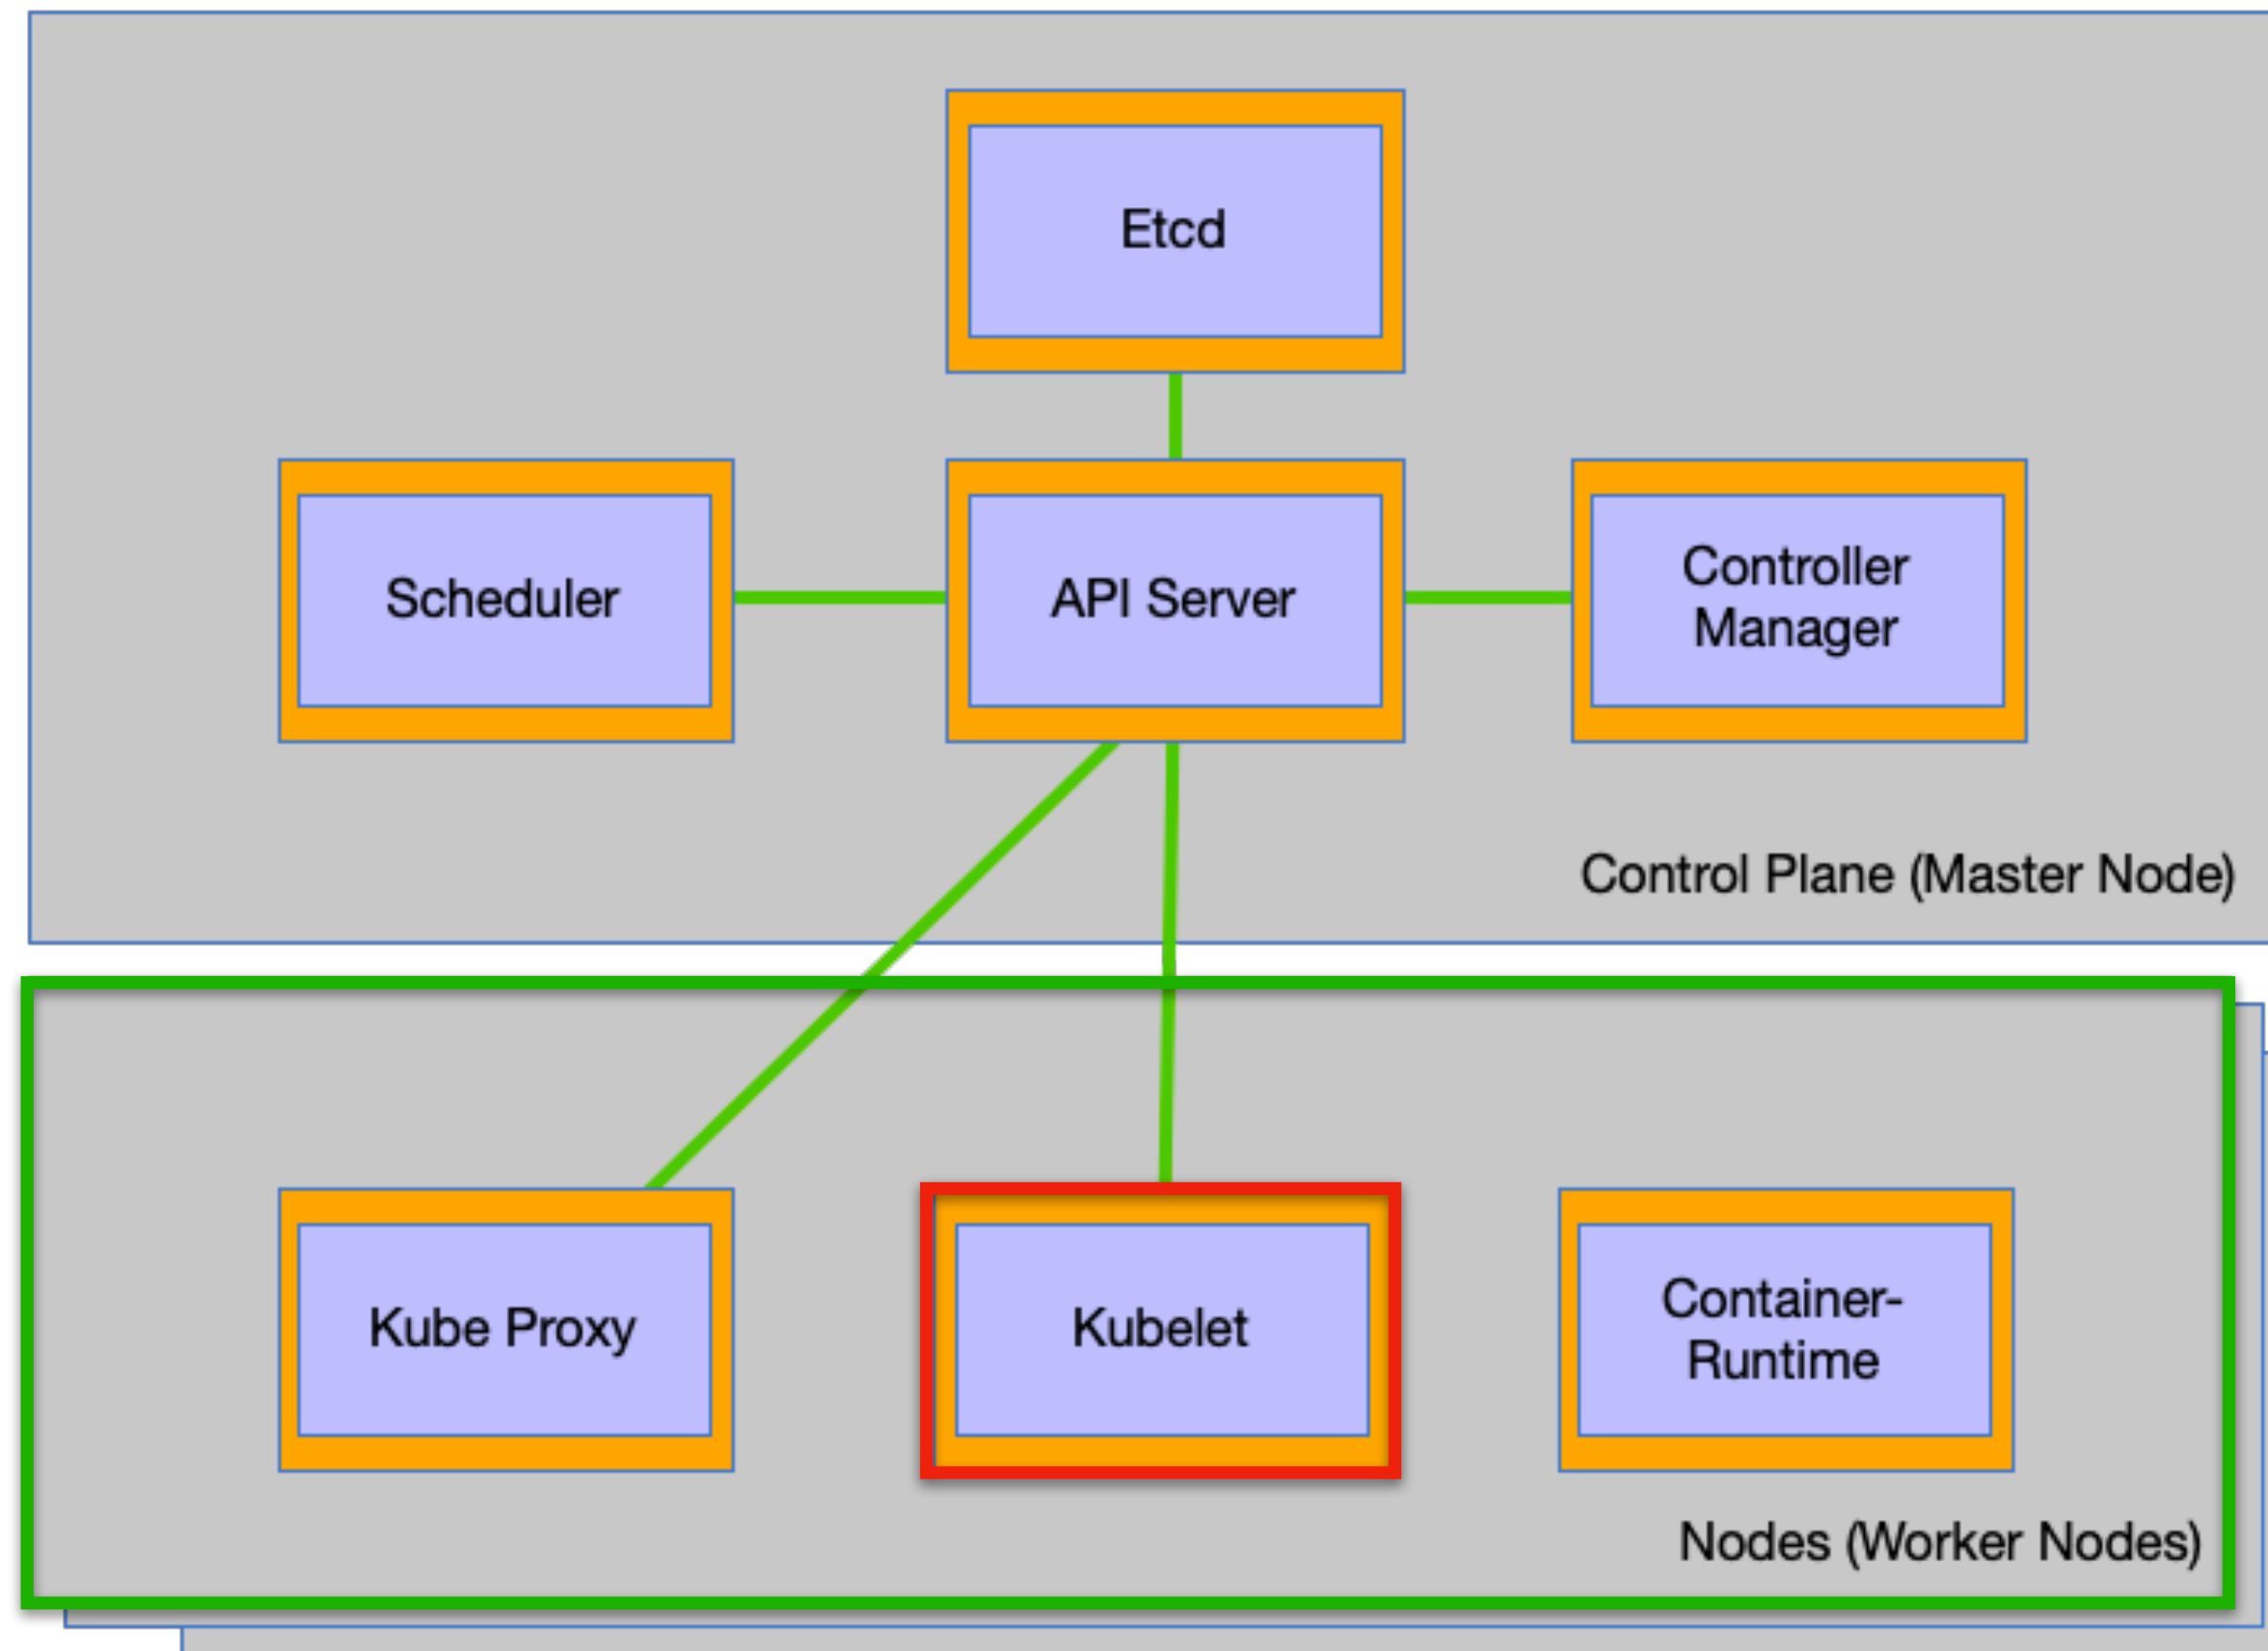
# Kube Proxy



- Network Proxy that runs in each node
- Just like performing <https://localhost>



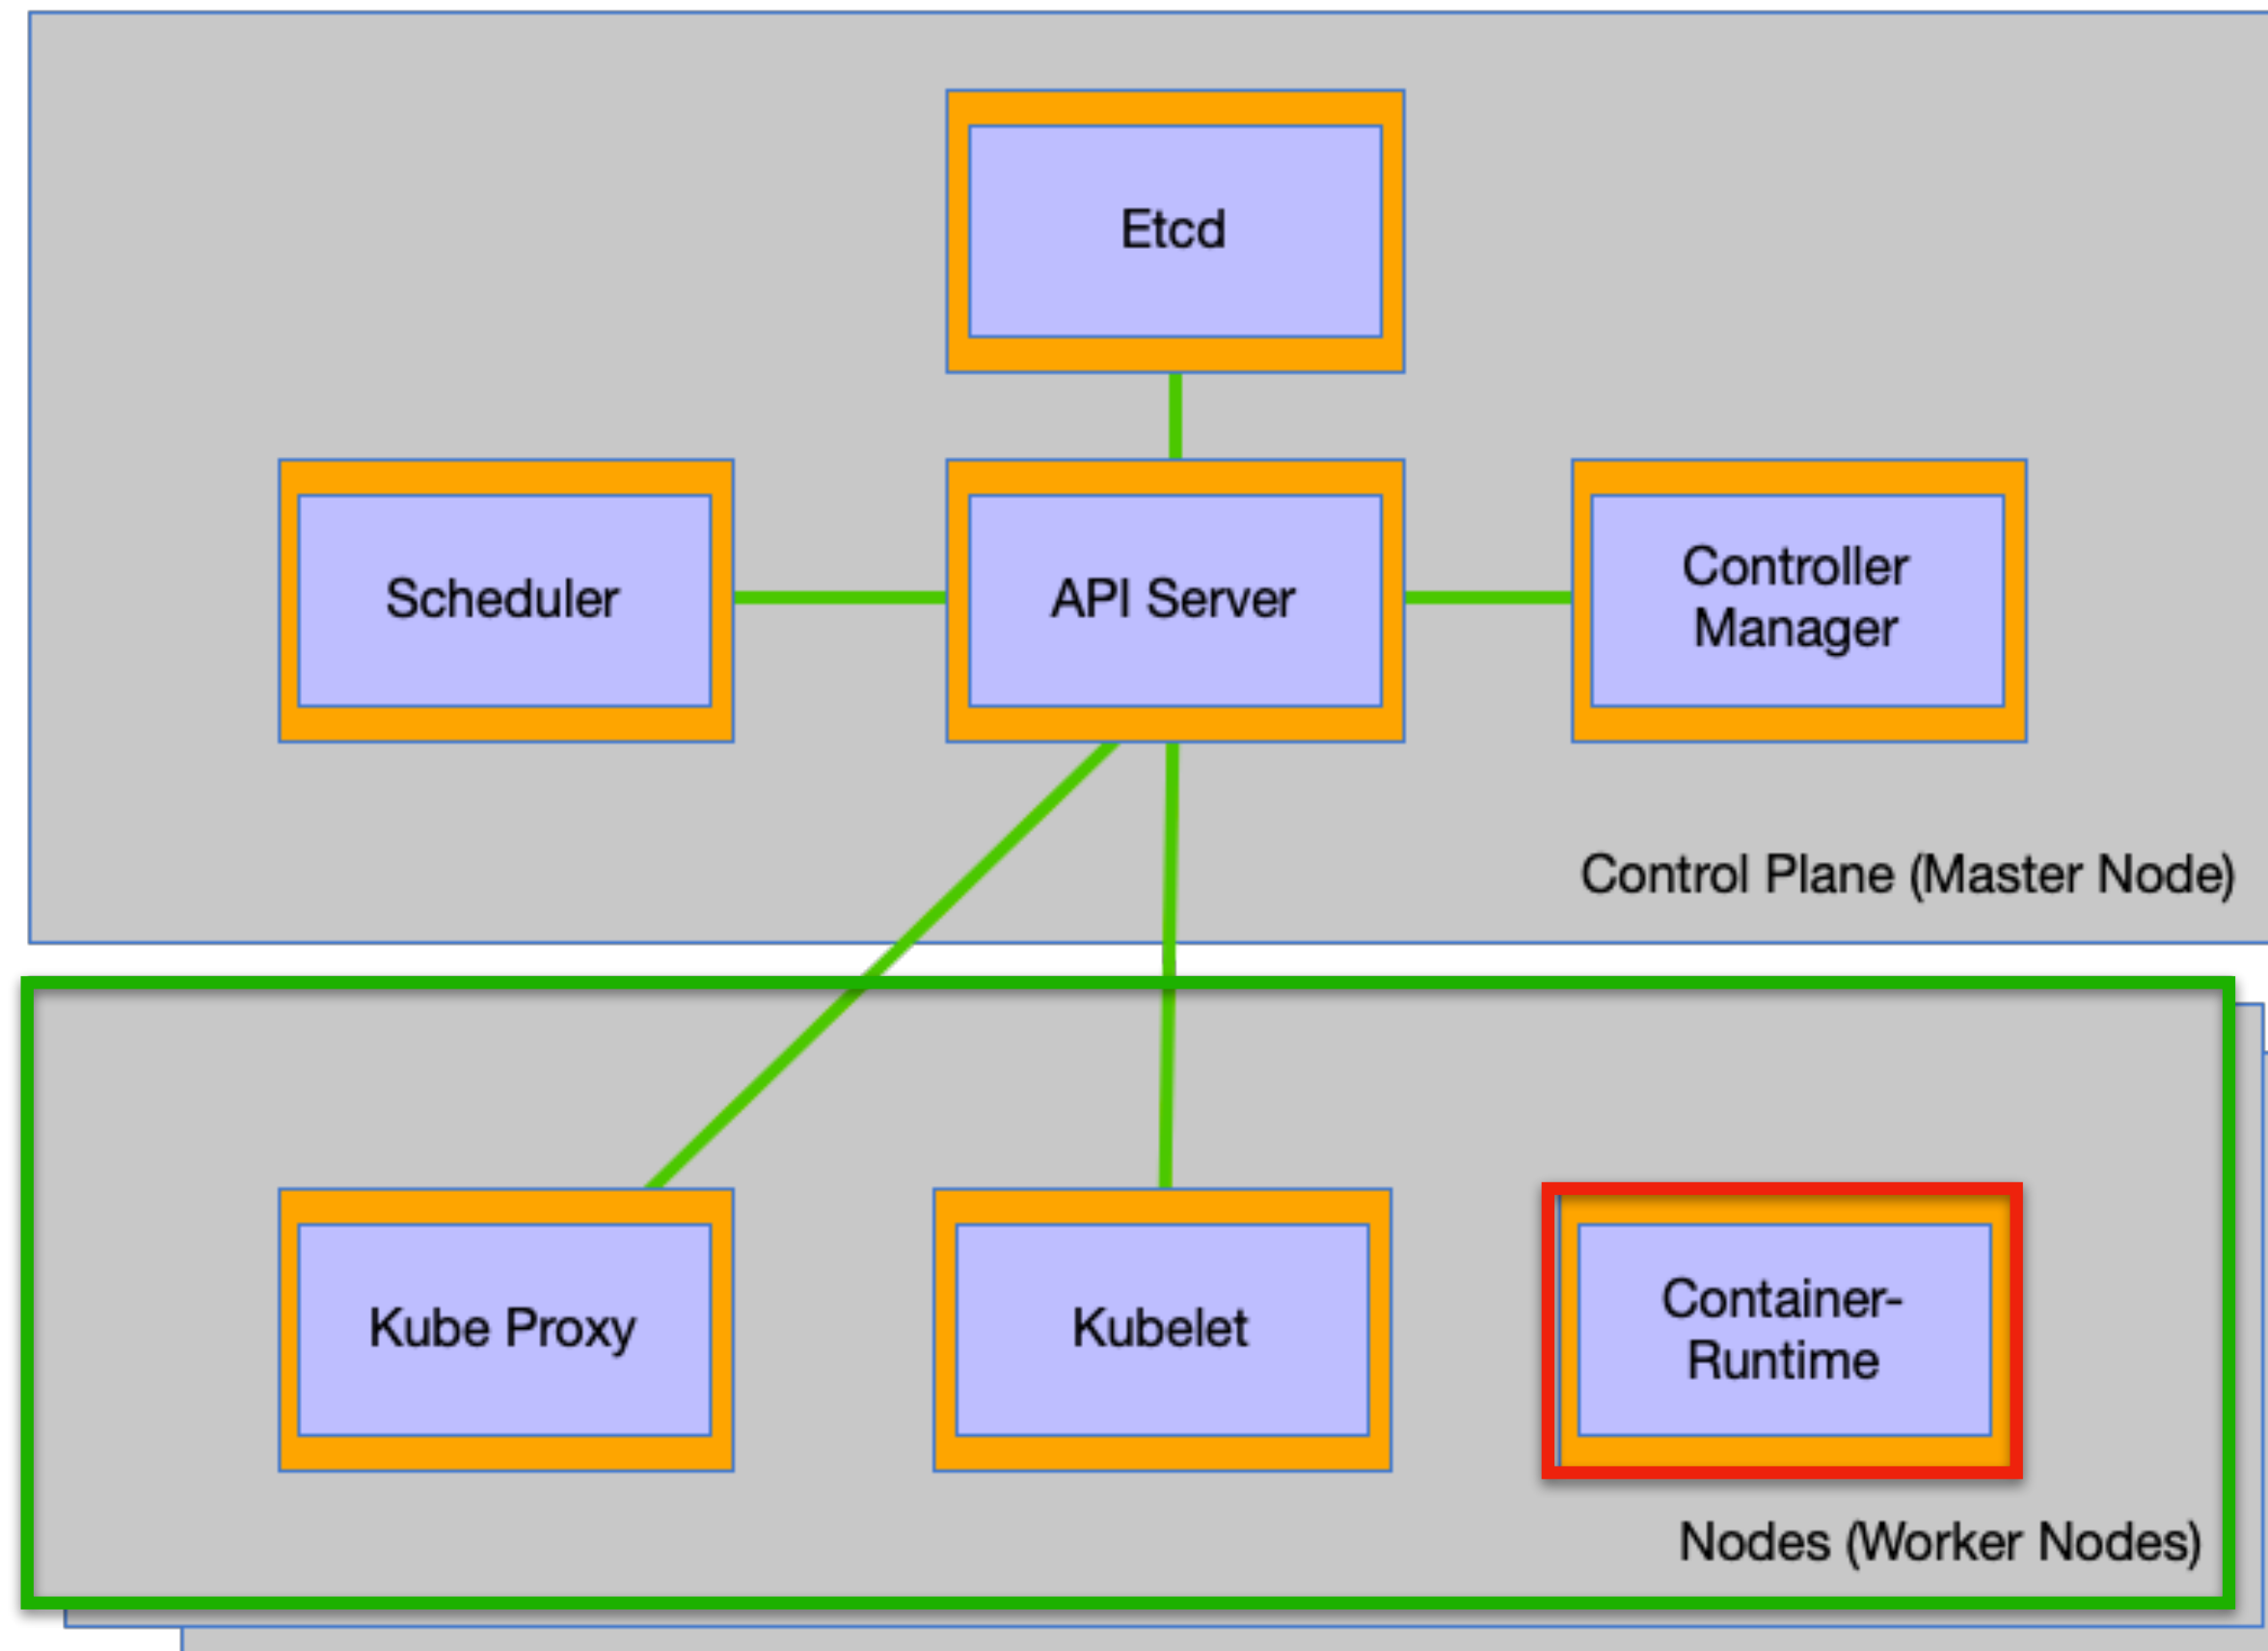
# Kubelet



- An agent that runs on each node in the cluster.
- Makes sure that containers are running in a Pod
- API Server connects to the Kubelet when fetching logs, attaching, or port-forwarding



# Container Runtime



- Software responsible for running the containers
- Supports varying container runtimes:
  - Docker
  - containerd
  - CRI-O
  - Any Implementation of the Kubernetes CRI (Container Runtime Interface)



# Understanding your Kube Config

The background of the slide features a blue sky with soft white clouds. Overlaid on this is a complex network of white lines connecting various sized white circular nodes, creating a mesh-like pattern that suggests a digital or networked environment.



# KubeConfig

- Default file: `.kube/config`
- File that maintains and organizes information about clusters, users, namespaces, and authentication mechanisms
- What is used by `kubectl` to find the information it needs to choose a cluster and communicate with the API server of a cluster
- Can be overridden with the `KUBECONFIG` environment variable
- Can use your own config file with `kubectl` using `--kubeconfig` flag
- Users and their certificates are also maintained within the `.kube/config`
- `kubectl config view` will show the active configuration

# Viewing KubeConfig Demo

Showing a KubeConfig Samples



# Kubernetes Objects Overview

The background of the slide features a blue sky with soft white clouds. Overlaid on this is a complex network of white lines connecting various sized white dots, creating a mesh-like pattern that suggests a digital or networked environment.



# TLS Termination Ingress

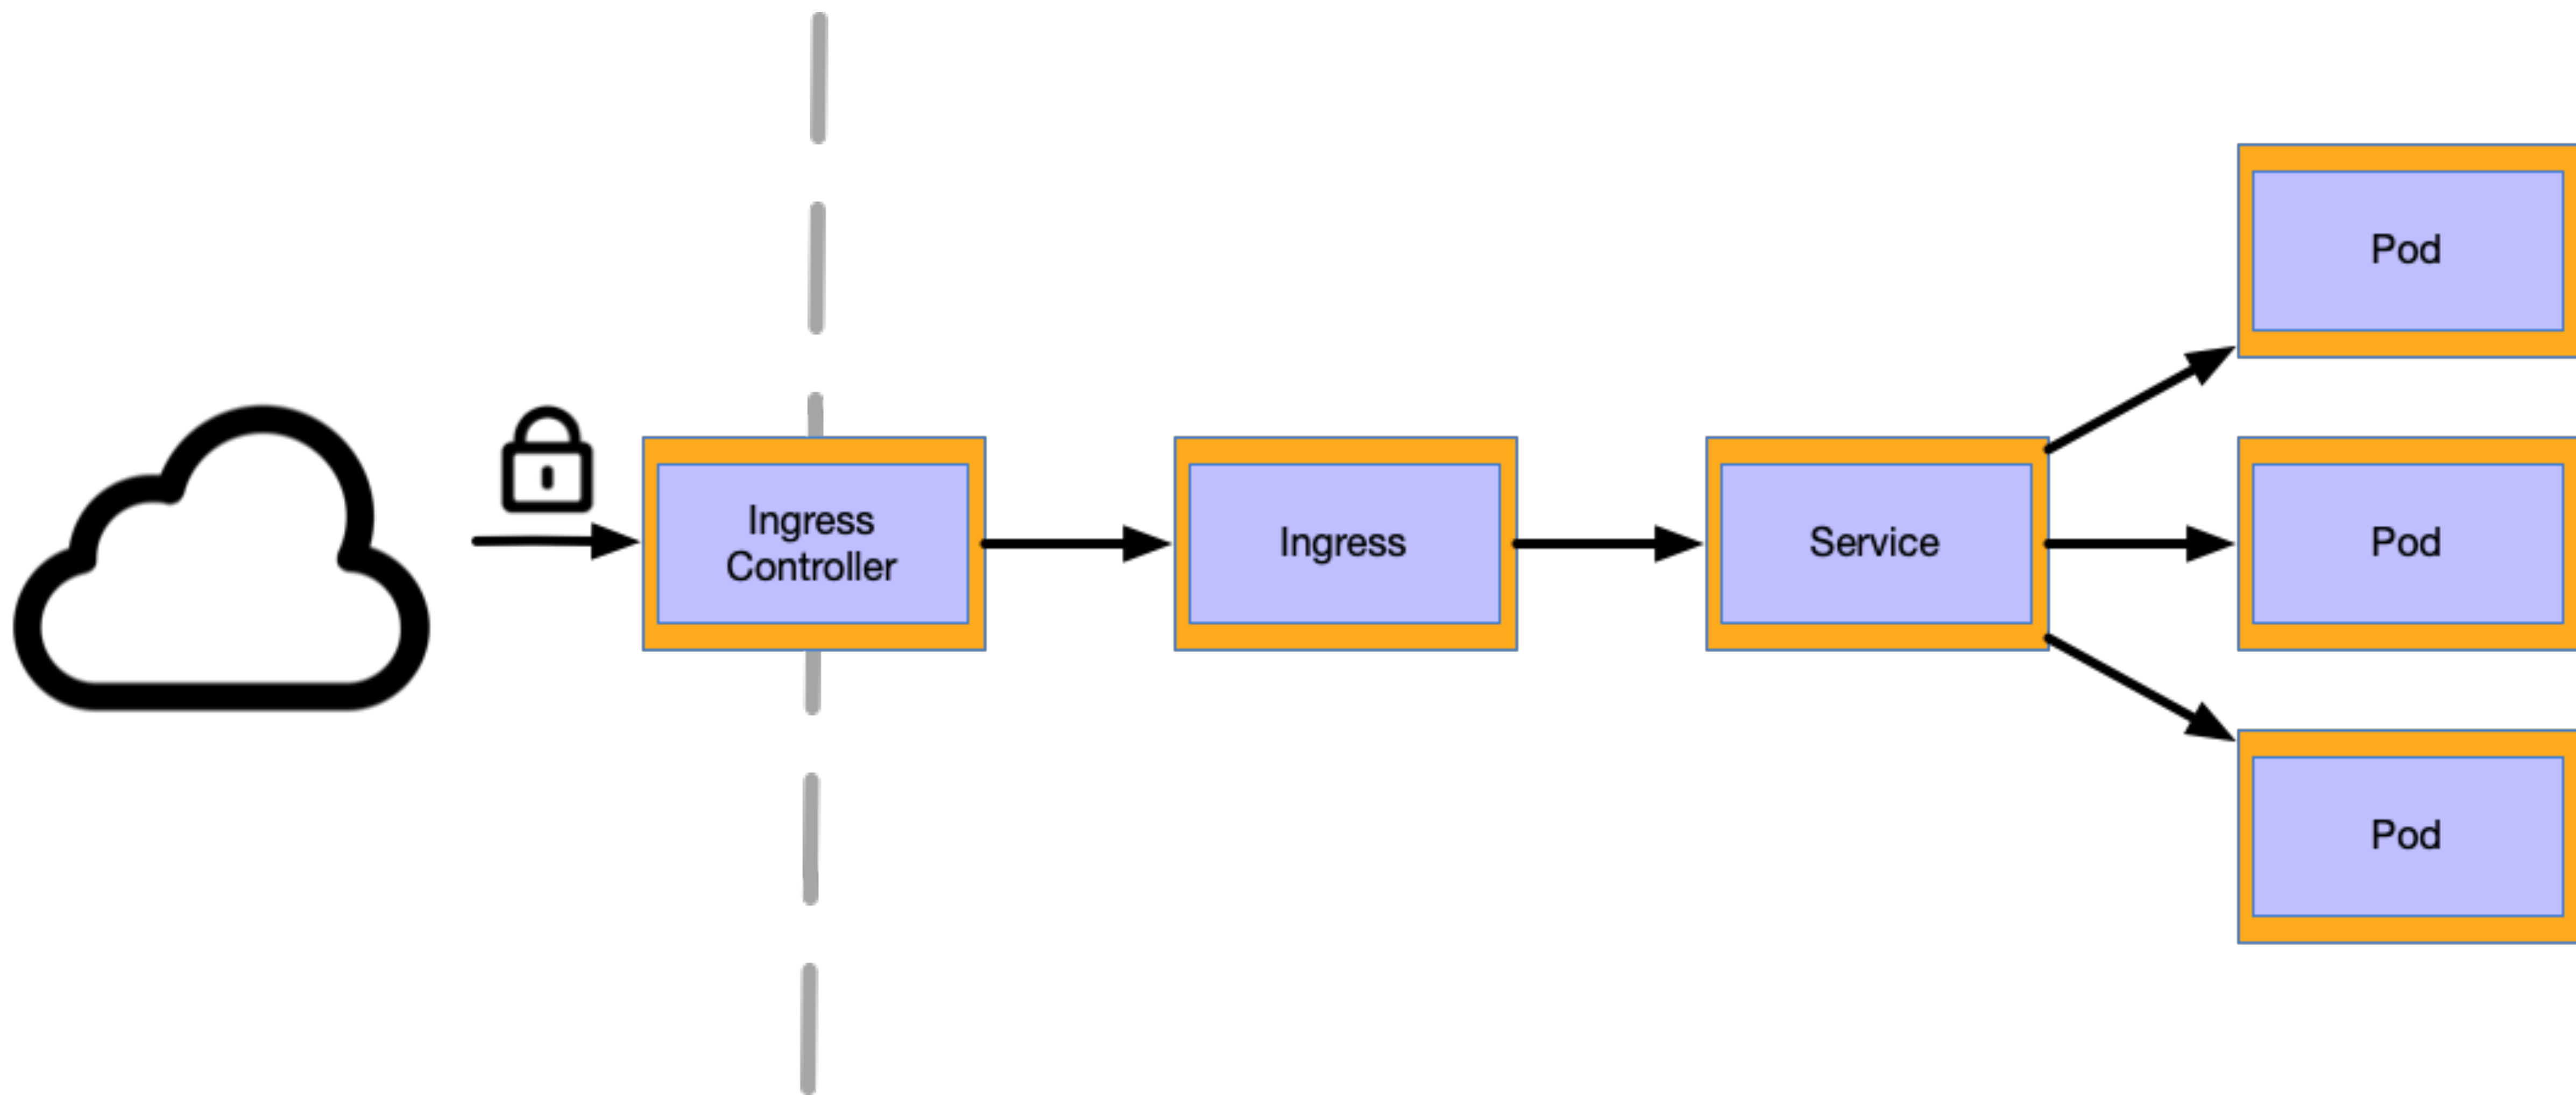
The background of the image features a blue sky with soft, white clouds. Overlaid on this is a complex network of white lines connecting various circular nodes. Some nodes are larger and more prominent, while others are smaller. The network structure is dense and spans across the lower half of the image, creating a sense of connectivity and data flow.



# Ingress Controllers

- Ingress Controller's a load balancing service that routes traffic to services within Kubernetes.
- Different Kubernetes environments use different implementations of the controller, but several don't provide a default controller at all
- A popular option would be to use nginx (perhaps via Helm)
- Ingress Controllers can also perform TLS-Termination, meaning that TLS traffic terminates at the controller

```
$ helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx  
$ helm install ingress-nginx ingress-nginx/ingress-nginx
```





# TLS Termination Ingress

Applying an Ingress and Certificate in AWS




# Kubernetes Authn & AuthZ

The background of the slide features a blue sky with soft white clouds. Overlaid on this is a complex network of white lines connecting various sized white dots, creating a mesh-like pattern that suggests a digital or networked environment.




# Least Privileged Security

- Minimum required privileges to perform a function
- Rewards of applying this principle
  - Better security
  - Better stability
- Improve audit readiness
- Authorization Models
  - **ACL**: list of permissions associated with objects
  - **RBAC**: subject's roles, contains group of permissions or privileges
  - **ABAC**: subject's attributes like labels or properties
- Recommended to enable RBAC in Kubernetes



- Client call attempt is made as an HTTP call to the API Server, this includes what we do with `kubectl`
- API Server is RESTful: `get`, `post`, `put`, `patch`, etc.
- The API Authenticates, Authorizes, Admits, and Validates before changing etcd






- Authentication is a list of the Authenticated Plugins that identifies who is gaining access to the Kubernetes API
- Authentication Plugins works like a linked list where one will be able to authenticate
- Once established then Authorization Plugins are negotiated

# Authentication Plugins

- X509 Client Certs
- Static Token File
- Bearer Tokens
- Bootstrap Tokens
- Service Account Tokens
- Open ID Connect Tokens
- Webhook Token Authentication





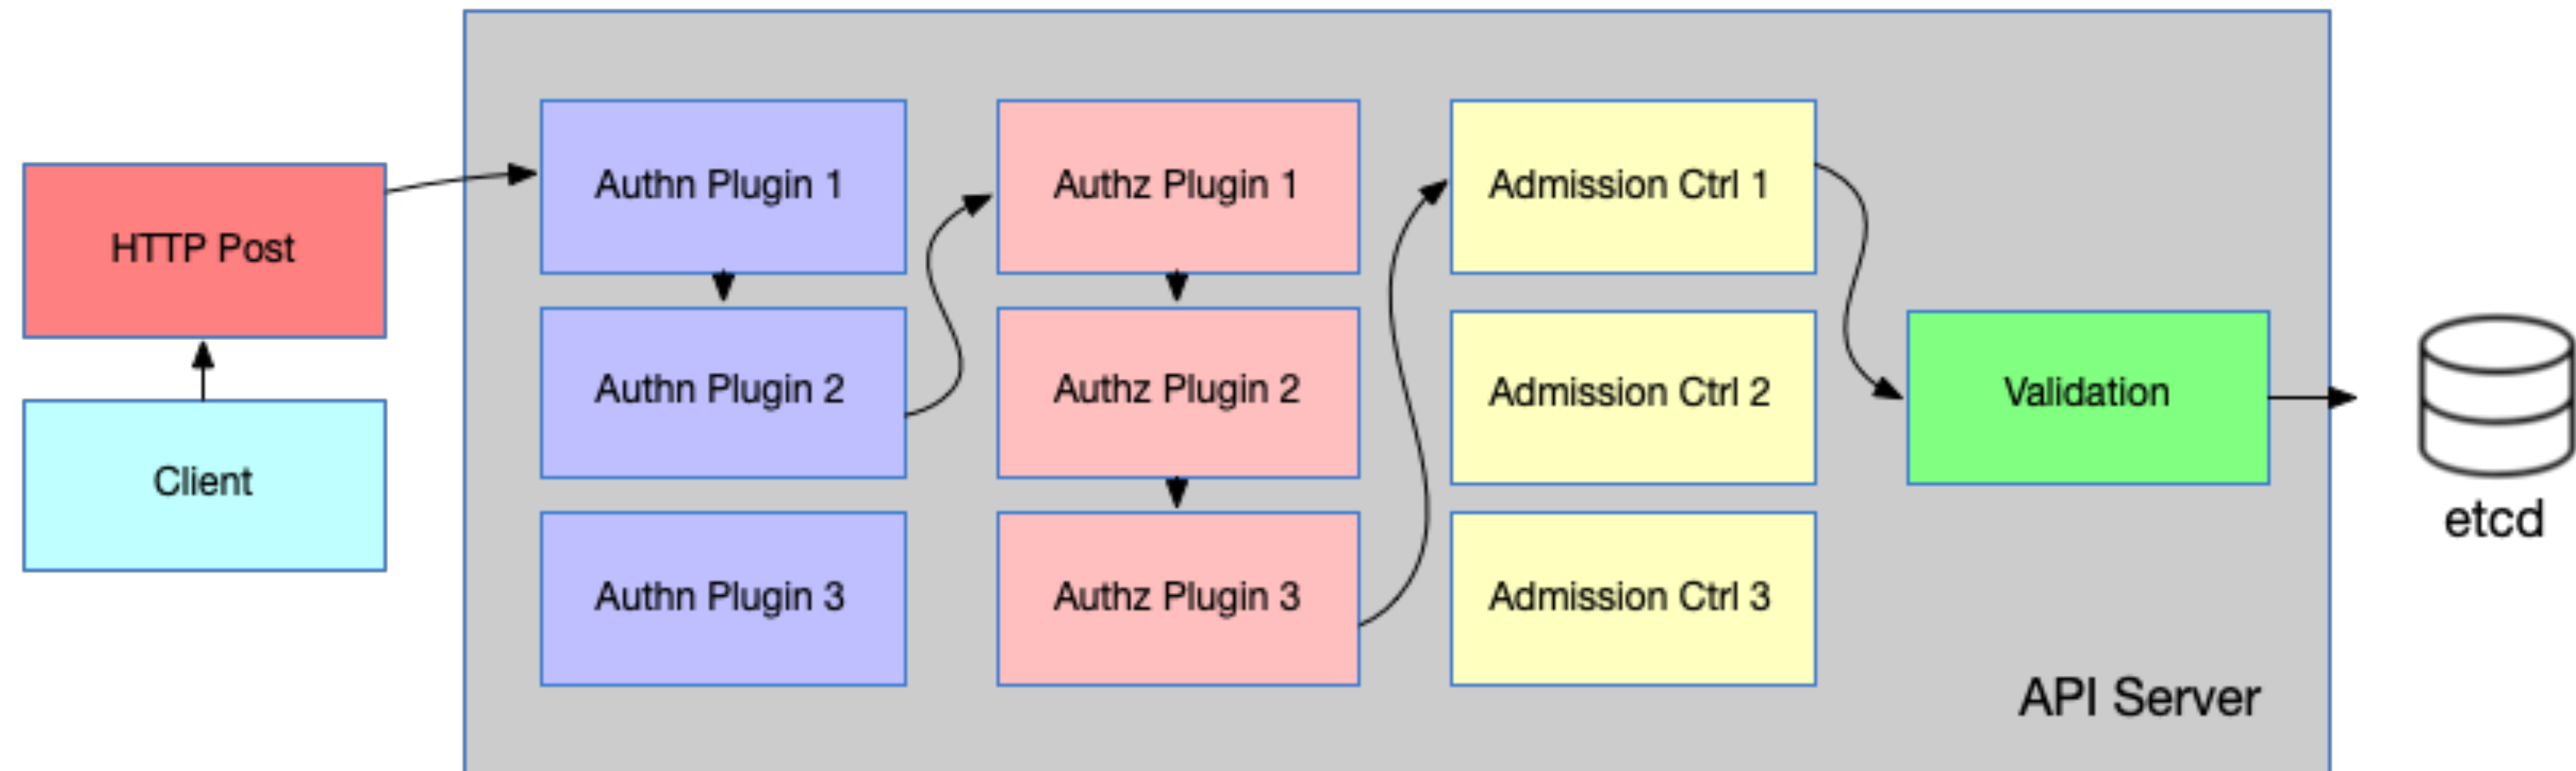
- Authorization Plugins works like a linked list where they have plugins
- Given the action that the user wishes to perform, the plugins will determine if the user is allowed to do so.
- As soon as a plugin says the user can perform the action, the API server progresses to the next stage



# Authorization Plugins

- RBAC Plugin - Role Based Access Control, checks whether an action is allowed to be performed by the user requesting the action
- ABAC Plugin - Attribute Based Access Control, defines an access control paradigm whereby access rights are granted to users through the use of policies which combine attributes together
- Node Plugin - Node authorization is a special-purpose authorization mode that specifically authorizes API requests made by kubelets
- Webhook Plugin - WebHook is an HTTP callback mode that allows you to manage authorization using a remote REST endpoint






- Admission Control plugins can modify the resource for different reasons
- They may initialize fields missing from the resource specification to the configured default values or even override them
- They may even modify other related resources, which aren't in the request, and can also reject a request for whatever reason
- GET (Read) calls don't go through the admission control plugins

# Admission Control Plugins

- `AlwaysPullImages`—Overrides the pod's `imagePullPolicy` to `Always`, forcing the image to be pulled every time the pod is deployed.
- `ServiceAccount`—Applies the default service account to pods that don't specify it explicitly.
- `NamespaceLifecycle`—Prevents creation of pods in namespaces that are in the process of being deleted, as well as in non-existing namespaces.
- `ResourceQuota`—Ensures pods in a certain namespace only use as much CPU and memory as has been allotted to the namespace.





- API server then validates the object, stores it in etcd, and returns a response to the client.
- Validation ensures that the data structure is correct before placing into etcd



# Service Accounts





# Kubernetes Types Of Accounts

- Humans
  - Meant to be Authenticated by External System
- Pods
  - Meant to be Authenticated by Service Account
  - Stored in the API Server as a Resource

# Service Accounts

- API Server authorizes requests coming from a pod based on a Service Account
- All pods by default are associated with the default Service Account
- Service Accounts
  - Can be associated with a namespace or cluster
  - Bound to Roles with a Role Binding
- Token Stored in each pod:  
    `/var/run/secrets/kubernetes.io/serviceaccount/token`



# Service Account Demo

Gaining Access to the Kubernetes API using Service Accounts



# User and Groups





# Users aren't your Standard Users

- There is no such user database in Kubernetes
- Users can be managed outside of Kubernetes (e.g. LDAP, SSO)
- No API Calls to add to Users
- Groups are also not defined in Kubernetes this is established as an outside concern

# Kubernetes Groups

- Users and Service Accounts belong to one or more Groups
- Authentication Plugins returns groups with username and user ID
- Groups are used to grant permission to many people at once
- Kubernetes has automatically created group that are returned with the following formats:
  - `system:unauthenticated` - unauthenticated clients
  - `system:authenticated` - user authenticated successfully
  - `system:serviceaccounts` - all service accounts
  - `system:serviceaccounts:<namespace>` - service accounts in a namespace



# Creating Users & Groups Demo

Establish a new user to access your Kubernetes Cluster



# RBAC





# RBAC

- Role Based Access Control
- General Availability as of Kubernetes 1.8
- Prevent unauthorized users from viewing or modifying the cluster state via the Kubernetes API
- Uses user roles as the key factor in determining whether the user may perform an action
- *A subject* (users, service account, or groups thereof) is associated with one or more roles
- A role is allowed to perform certain verbs (GET, POST, PUT, PATCH, DELETE)

# Http Methods and Verbs

What do they mean when consulting the API

HTTP Method	Verb Single	Verb Plural
GET, HEAD	get (and watch)	list (and watch)
POST	create	
PUT	update	
PATCH	patch	
DELETE	delete	deletecollection



# Role Bindings





# Role

- Always sets permissions within a particular namespace
- Must add a namespace that it belongs
- Roles define what can be done



# Defining a Role

## Establishing a Role within Kubernetes

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: data-engineering
  name: pod-reader
rules:
- apiGroups: [""] # Core API
  verbs: ["get", "list"]
  resources: ["pods"]
```

# Defining a Role Binding to Service Account

## Establishing a RoleBinding within Kubernetes

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: test
  namespace: foo
  ...
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: service-reader
subjects:
- kind: ServiceAccount
  name: default
  namespace: foo
```



# Role Binding Demo

Allow Users and Service Account a little more free reign within your Kubernetes Cluster



# Cluster Bindings





# ClusterRole

- Role that represents a non-namespaced resources
- Resources that are not namespaced includes Nodes, PersistentVolumes, and Namespaces
- The Kubernetes API can expose some URLs that don't represent resources, like the `/healthz` endpoint
- Namespaced resources (like Pods), across all namespaces For example: you can use a ClusterRole to allow a particular user to run `kubectl get pods --all-namespaces`

# ClusterRole YAML Example

ClusterRoles have no namespaces

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  # "namespace" omitted since ClusterRoles are not namespaced
  name: secret-reader
rules:
- apiGroups: [""]
  #
  # at the HTTP level, the name of the resource for accessing Secret
  # objects is "secrets"
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
```



# ClusterRoleBindings

- Binds together the cluster role with subject (User, Group, ServiceAccount)
- Kubernetes comes with a default set of ClusterRoles and ClusterRoleBindings
- Cluster Role Bindings Updated everytime the API starts

# ClusterRoleBinding YAML Example

## ClusterRoleBinding to either a User, Group, or ServiceAccount

```
apiVersion: rbac.authorization.k8s.io/v1
# This cluster role binding allows anyone in the "manager" group to
# read secrets in any namespace.
kind: ClusterRoleBinding
metadata:
  name: read-secrets-global
subjects:
- kind: Group
  name: manager # Name is case sensitive
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: secret-reader
  apiGroup: rbac.authorization.k8s.io
```



# Cluster Role Binding Demo

Give a user access to do more to your Kubernetes Cluster



# Securing Pods

The background of the slide is a vibrant blue sky filled with soft, white clouds. Overlaid on this is a complex network of white lines and dots, resembling a digital or neural network. The dots vary in size, and the lines connect them in a web-like pattern, creating a sense of connectivity and data flow. The overall aesthetic is modern and technological.



# Networking Policies

The background of the slide features a blue sky with white clouds. Overlaid on this is a complex network diagram. The diagram consists of numerous white dots of varying sizes, representing nodes, which are interconnected by thin white lines, representing network connections. The network is dense and spans the entire width of the slide, with some areas appearing more concentrated than others.



# Network Policies


- NetworkPolicy applies to pods that match its label selector
- Specifies either which sources can access the matched pods or which destinations can be accessed from the matched pods
- Operated with ingress or egress **rules**. Ingress here has nothing to do with the ingress as we talked about previously
- Matching pods can be performed using:
  - A pod selector
  - A namespace selector
  - CIDR Notation (192.168.1.1/24)



# Network Policy by Pod Selector

Selecting the traffic allowed inwards, ingress


```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: database-netpolicy
spec:
  podSelector:
    matchLabels:
      app: database
  ingress:
    - from:
      - podSelector:
          matchLabels:
            app: webserver
    ports:
      - port: 3066
```



# Network Policy by Pod Selector

Selecting the traffic allowed outwards, egress

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: webserver-netpolicy
spec:
  podSelector:
    matchLabels:
      app: webserver
  egress:
    - from:
      - podSelector:
          matchLabels:
            app: database
```





# Essential Tools





# KubeBench

- <https://github.com/aquasecurity/kube-bench>
- kube-bench is a tool that checks whether Kubernetes is deployed securely by running the checks documented in the CIS Kubernetes Benchmark

```
$ kubectl apply -f job.yaml
job.batch/kube-bench created

$ kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
kube-bench-j76s9    0/1     ContainerCreating   0           3s

# Wait for a few seconds for the job to complete
$ kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
kube-bench-j76s9    0/1     Completed      0          11s

# The results are held in the pod's logs
kubectl logs kube-bench-j76s9
[INFO] 1 Master Node Security Configuration
[INFO] 1.1 API Server
...
```



# Tracee

- Observe system calls made by your container
- Uses eBPF as a sandbox and determine what capabilities your container uses
- Once the report is done, you can remove capabilities from image

```
docker run \  
  --name tracee --rm -it \  
  --pid=host --cgroupns=host --privileged \  
  -v /etc/os-release:/etc/os-release-host:ro \  
  -e LIBBPF_GO_OSRELEASE_FILE=/etc/os-release-host \  
  aquasec/tracee:latest
```

```
docker run -it --rm -p 8080:80 nginx:alpine
```

```
docker run --cap-drop=all --cap-add=<cap1> --cap-add=<cap2> <image> ...
```

# Open Policy Agent

- Admission Controllers enforce semantic validation of objects during create, update, and delete operations.
- With OPA you can enforce custom policies on Kubernetes objects without recompiling or reconfiguring the Kubernetes API server. Uses eBPF as a sandbox and determine what capabilities your container uses
- Uses a configuration language called Rego to enforce those policies
- <https://www.openpolicyagent.org/docs/v0.12.2/kubernetes-admission-control/>



Open Policy Agent



# Trivy

- Trivy has different *scanners* that look for different security issues, and different *targets* where it can find those issues.
- Trivy's targets can include docker containers, file systems, and Kubernetes clusters
- <https://aquasecurity.github.io/trivy/v0.28.1/>

```
trivy image python:3.4-alpine
```

```
trivy fs --security-checks vuln,secret,config myproject/
```

```
trivy k8s mycluster
```

# Kube-Hunter

- kube-hunter hunts for security weaknesses in Kubernetes clusters [Pen Tool]
- The tool was developed to increase awareness and visibility for security issues in Kubernetes environments.
- Trivy's targets can include docker containers, file systems, and Kubernetes clusters
- <https://aquasecurity.github.io/trivy/v0.28.1/>

```
apiVersion: batch/v1
kind: Job
metadata:
  name: kube-hunter
spec:
  template:
    spec:
      containers:
      - name: kube-hunter
        image: aquasec/kube-hunter
        command: ["kube-hunter"]
        args: ["--pod"]
        restartPolicy: Never
      backoffLimit: 4
```



# Falco

- A CNCF incubating project
- Falco is an open source runtime security tool
- Falco parses Linux system calls from the kernel at runtime
- Asserts the stream against a rules engine
- Run as a Kubernetes Job
- If a rule is violated, Falco triggers an alarm

```
helm install falco falcosecurity/falco  
kubectl get pods
```



# Falco, what does it check by default?

- Privilege escalation using privileged containers
- Namespaces changes using tools like setns
- Read/Write to well-known directories such as /etc, /usr/bin
- Creating symlinks
- Ownership and mode changes
- Executing SSH or shell binaries



# Pod Security Admission

The background of the slide features a blue sky with soft white clouds. Overlaid on this is a complex network of white lines connecting various sized white circular nodes, creating a mesh-like pattern that suggests a digital or networked environment.



# Pod Security Admission

- Successor to Pod Security Policies which is now deprecated
- Used to enforce rules *per namespace* as to what kind of pods can be deployed with privilege levels: `privileged`, `baseline`, and `restricted`
- Kubernetes defines a set of labels that you can set to define which of the predefined Pod Security Standard levels you want to use for a namespace.
- The label you select defines what action the control plane takes if a potential violation is detected
- Exemptions for pods can be declared (e.g. users, class names, namespaces)

Mode	Description
<b>enforce</b>	Policy violations will cause the pod to be rejected.
<b>audit</b>	Policy violations will trigger the addition of an audit annotation to the event recorded in the <a href="#">audit log</a> , but are otherwise allowed.
<b>warn</b>	Policy violations will trigger a user-facing warning, but are otherwise allowed.



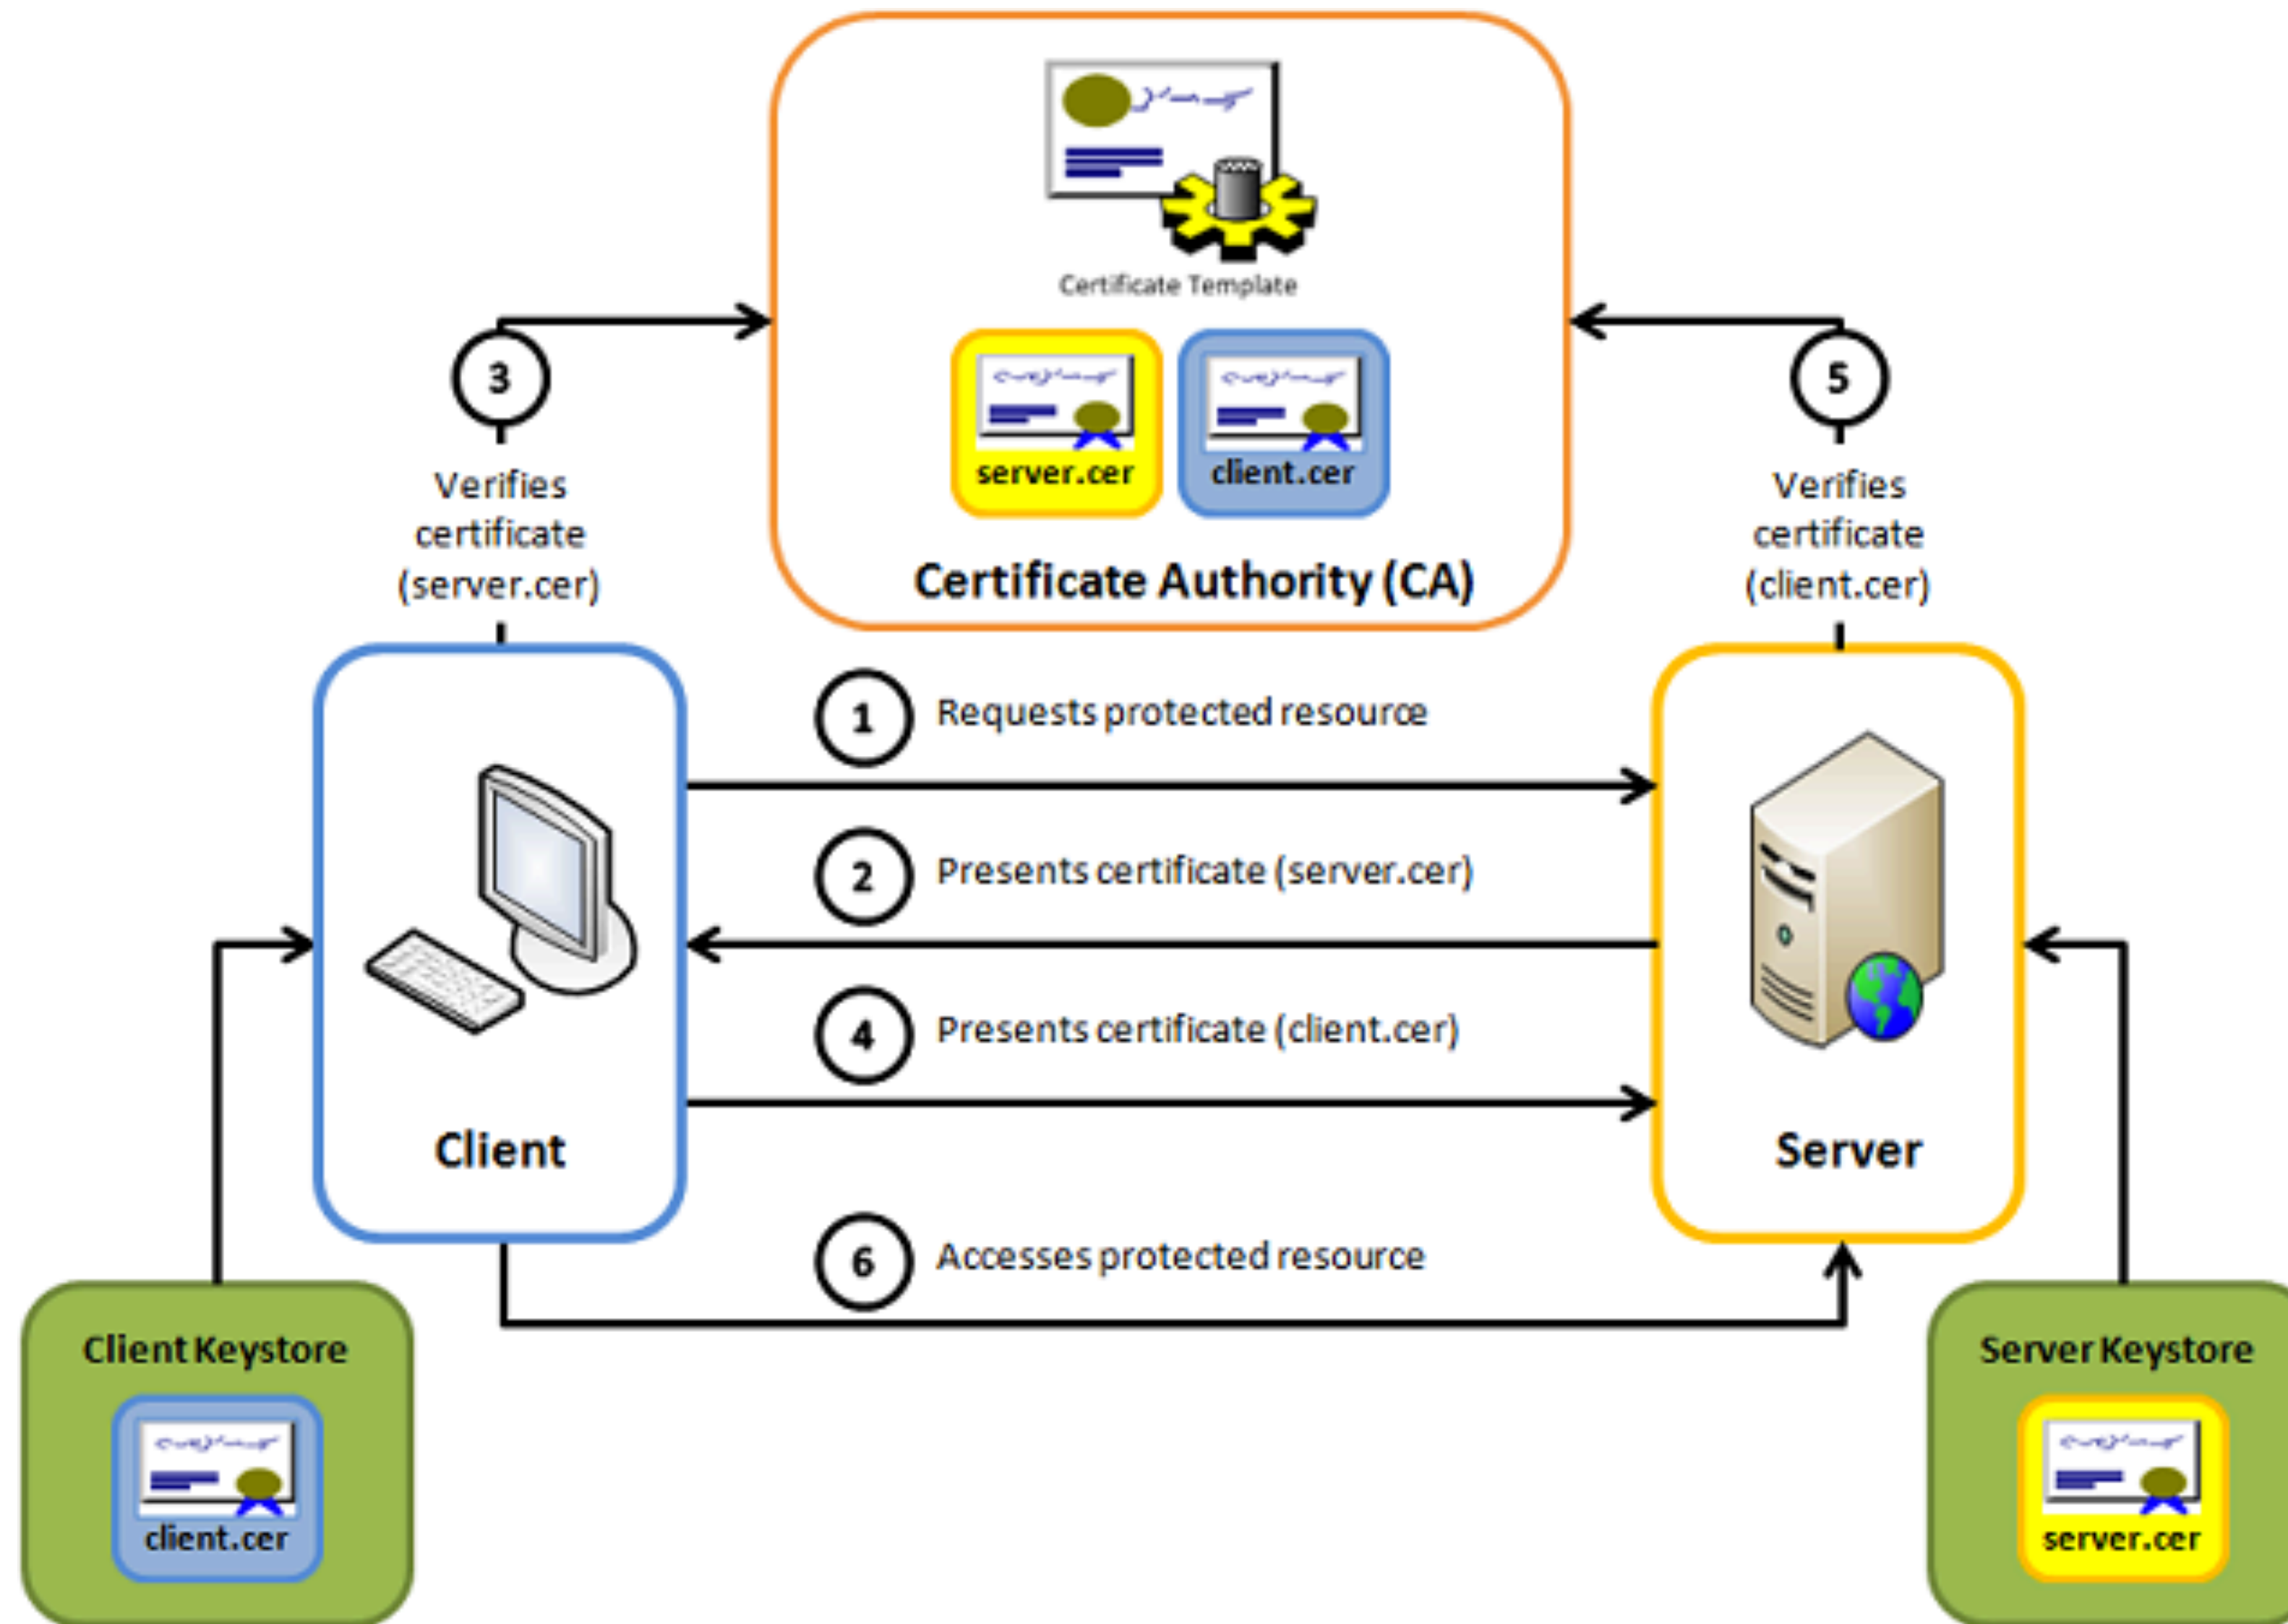
# Outer Network mTLS

The background of the slide features a blue sky with soft white clouds. Overlaid on this is a complex network diagram consisting of numerous white circular nodes of varying sizes, interconnected by thin white lines. The network is more densely packed in the lower half of the image, creating a sense of depth and connectivity.



# mTLS Defined

Mutual TLS (mTLS) - Authentication ensures that traffic is both secure and trusted **in both directions** between a client and server.



**Mutual SSL authentication / Certificate based mutual authentication**



# Step for mTLS

- Setup an Ingress Controller
- Create Certificates
- Create Kubernetes Secrets
- Deploy your Application

# Creating the Certificates

- **CommonName(CN):** Identifies the hostname or owner associated with the certificate.
- **Certificate Authority(CA):** A trusted 3rd party that issues Certificates. Usually you would obtain this from a trusted source, but for this example we will just create one. The CN is usually the name of the issuer.
- **Server Certificate:** A Certificate used to identify the server. The CN here is the hostname of the server. The Server Certificate is valid only if it is installed on a server where the hostname matches the CN.
- **Client Certificate:** A Certificate used to identify a client/user. The CN here is usually the name of the client/user.



# Creating the mTLS Certificates

```
# Generate the CA Key and Certificate
```

```
$ openssl req -x509 -sha256 -newkey rsa:4096 -keyout ca.key -out ca.crt -days 356 -nodes  
-subj '/CN=Fern Cert Authority'
```

```
# Generate the Server Key, and Certificate and Sign with the CA Certificate
```

```
$ openssl req -new -newkey rsa:4096 -keyout server.key -out server.csr -nodes -subj '/  
CN=meow.com'
```

```
$ openssl x509 -req -sha256 -days 365 -in server.csr -CA ca.crt -CAkey ca.key -set_serial  
01 -out server.crt
```

```
# Generate the Client Key, and Certificate and Sign with the CA Certificate
```

```
$ openssl req -new -newkey rsa:4096 -keyout client.key -out client.csr -nodes -subj '/  
CN=Fern'
```

```
$ openssl x509 -req -sha256 -days 365 -in client.csr -CA ca.crt -CAkey ca.key -set_serial  
02 -out client.crt
```

# Storing CA and Server CRT

```
$ kubectl create secret generic my-certs --from-file=tls.crt=server.crt --from-file=tls.key=server.key --from-file=ca.crt=ca.crt
```

```
$ kubectl get secret my-certs
```

NAME	TYPE	DATA	AGE
my-certs	Opaque	3	1m

- Deploy your application as normal with Services and Deployments
- We will configure the Ingress to use my-certs as to where the certificate authority and server certificates will reside



# Establishing Ingress Certificates

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/auth-tls-verify-client: \"on\"
    nginx.ingress.kubernetes.io/auth-tls-secret: \"default/my-certs\"
  name: meow-ingress
  namespace: default
```

- TLS is enabled and it is using the `tls.key` and `tls.crt` provided in the `my-certs` secret.
- The `nginx.ingress.kubernetes.io/auth-tls-secret` annotation uses `ca.crt` from the `my-certs` secret.

# Testing mTLS

```
$ curl https://meow.com/ -k
...
<center><h1>400 Bad Request</h1></center>
<center>No required SSL certificate was sent</center>
....
```

- -k is insecure, don't consult with a certificate verification
- In the following, a client certification and client key is used to get the payload

```
$ curl https://meow.com/ --cert client.crt --key client.key -k
...
ssl-client-issuer-dn=CN=Fern Cert Authority
ssl-client-subject-dn=CN=Fern
ssl-client-verify=SUCCESS
user-agent=curl/7.54.0
...
```



# Inner Network TLS





# TLS The Hard Way

```
cat <<EOF | cfssl genkey - | cfssljson -bare server
{
  "hosts": [
    "my-svc.my-namespace.svc.cluster.local",
    "my-pod.my-namespace.pod.cluster.local",
    "192.0.2.24",
    "10.0.34.2"
  ],
  "CN": "system:node:my-pod.my-namespace.pod.cluster.local",
  "key": {
    "algo": "ecdsa",
    "size": 256
  },
  "names": [
    {
      "0": "system:nodes"
    }
  ]
}
EOF
```



# Create Certificate Signing Request

```
cat <<EOF | kubectl apply -f -
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: my-svc.my-namespace
spec:
  request: $(cat server.csr | base64 | tr -d '\n')
  signerName: kubernetes.io/kubelet-serving
  usages:
    - digital signature
    - key encipherment
    - server auth
EOF
```

# Creating the Secret

```
apiVersion: "v1"
kind: "Secret"
metadata:
  name: "nginxsecret"
  namespace: "default"
type: kubernetes.io/tls
data:
  tls.crt: "LS0tL..."
  tls.key: "LS0tL..."
```

- Apply the TLS.crt and TLS.key as a Kubernetes Secret
- This can be used when defining the Service



# Creating the Secret

```
spec:
  volumes:
    - name: secret-volume
      secret:
        secretName: nginxsecret
    ...
  containers:
    - name: nginxhttps
      image: bprashanth/nginxhttps:1.0
      ports:
        ...
      volumeMounts:
        - mountPath: /etc/nginx/ssl
          name: secret-volume
```

- Create a Secret Volume which refers to the Secret established in the previous slide
- Bind the SSL to the directory where in this case NGINX is requiring the certificates



# Istio/Service Mesh and mTLS

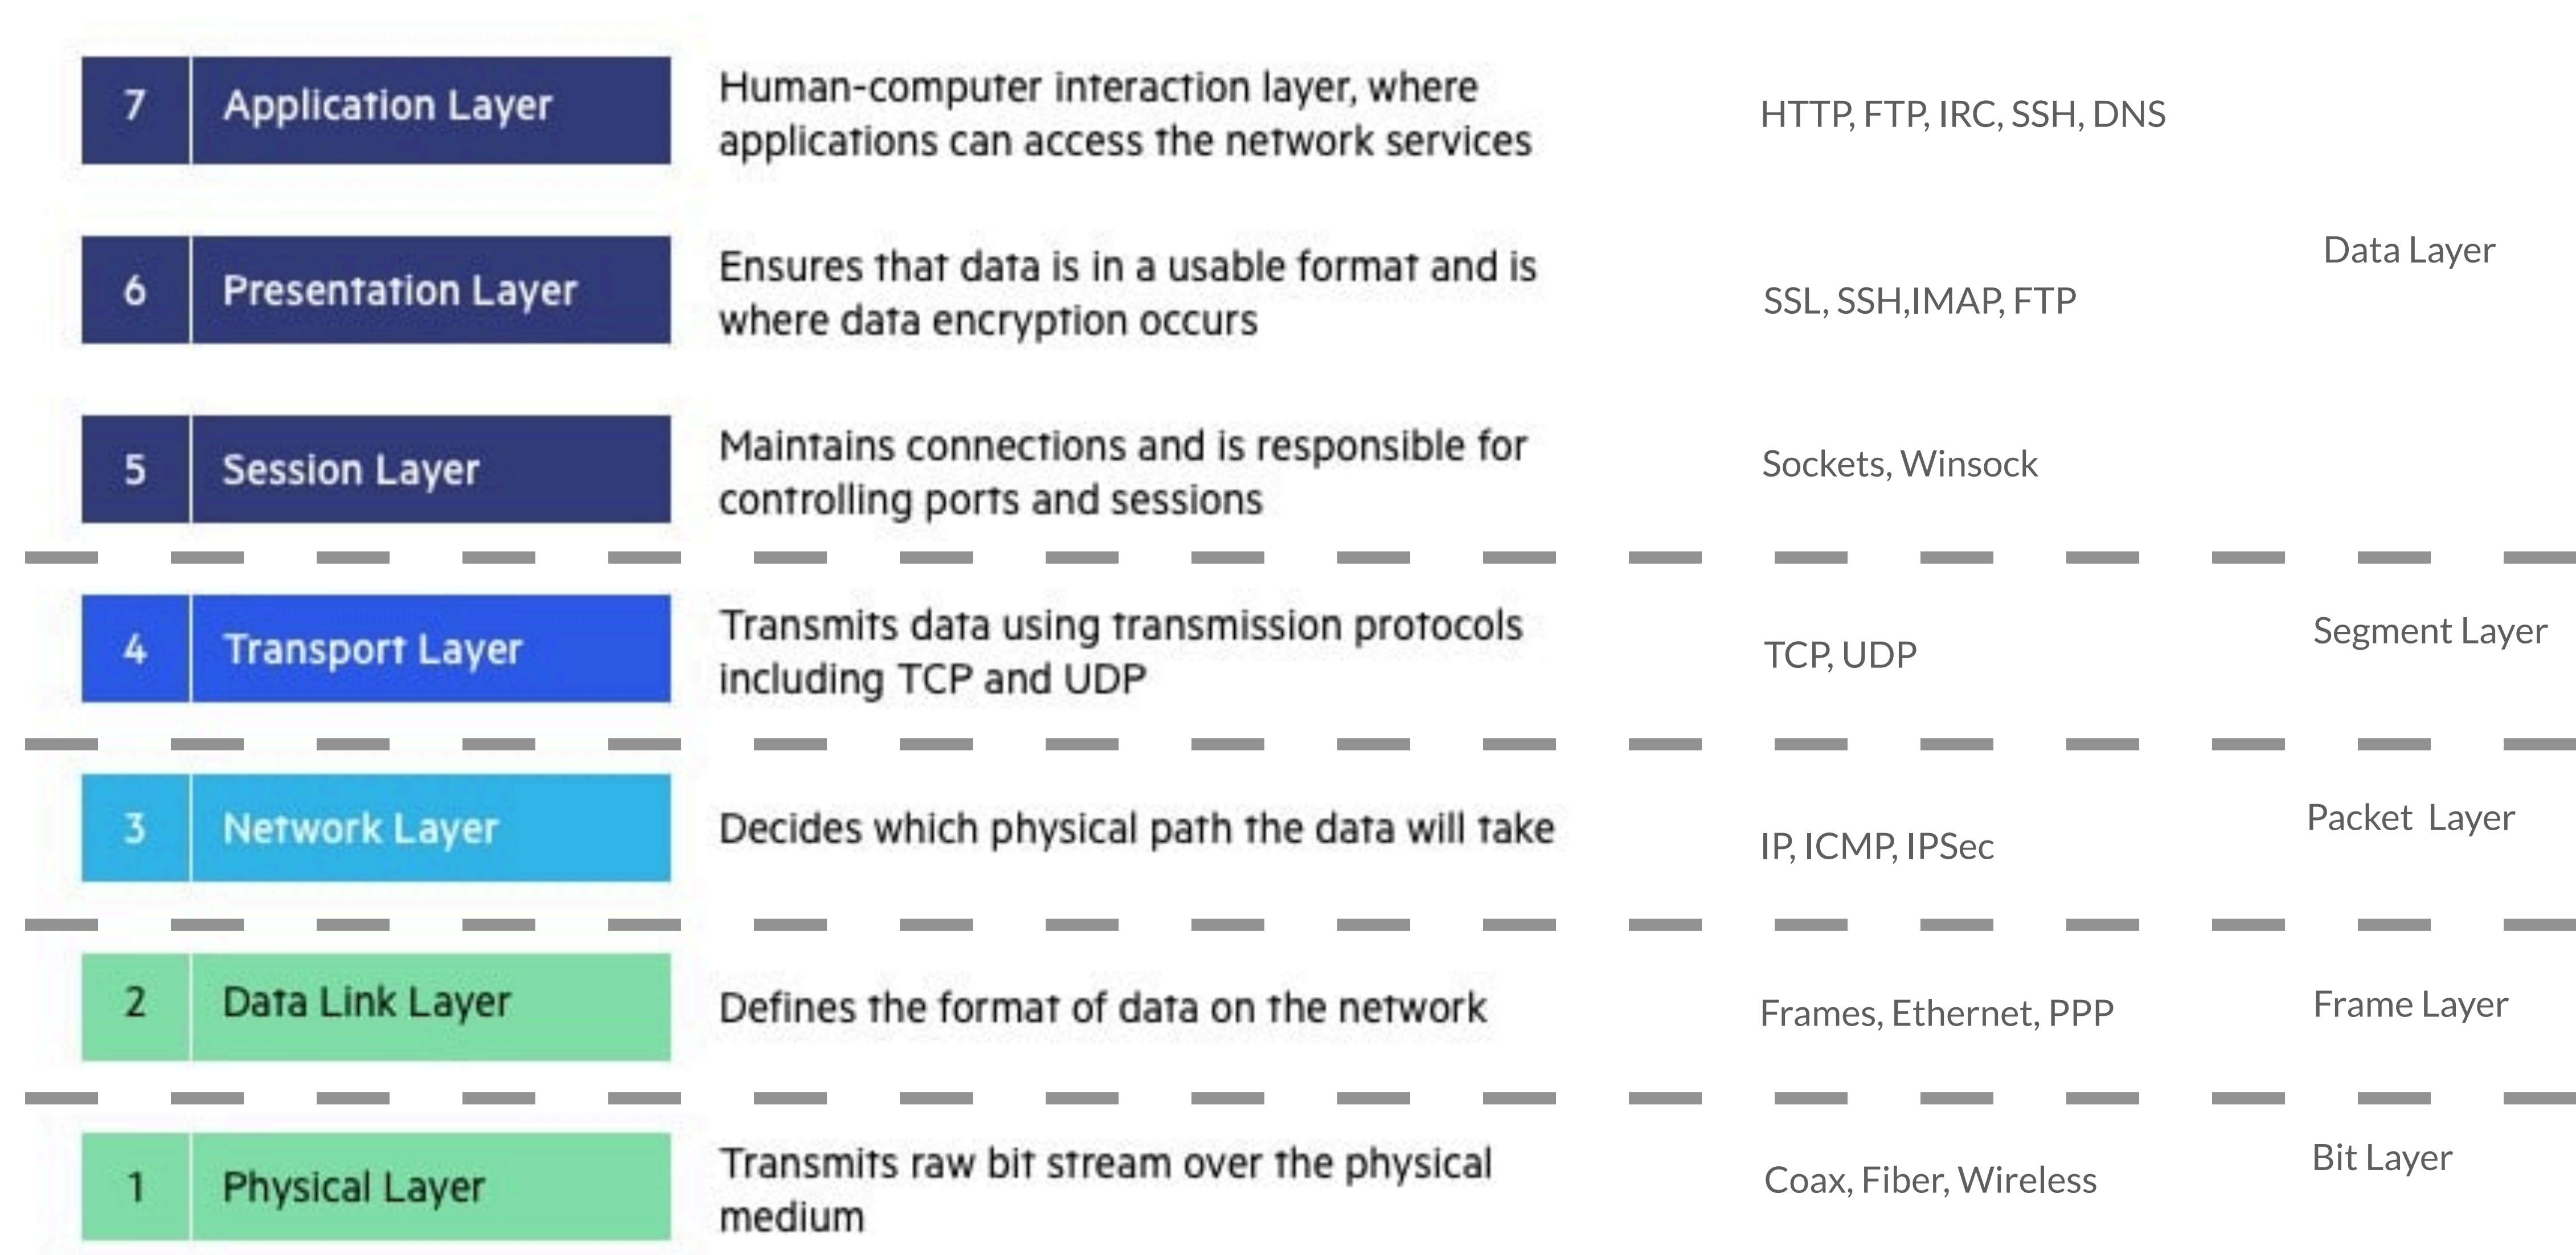
The background of the slide features a blue sky with soft white clouds. Overlaid on this is a complex network diagram consisting of numerous white circular nodes of varying sizes, interconnected by thin white lines. The network is more densely packed in the lower half of the image, creating a sense of depth and connectivity.



# Service Meshes

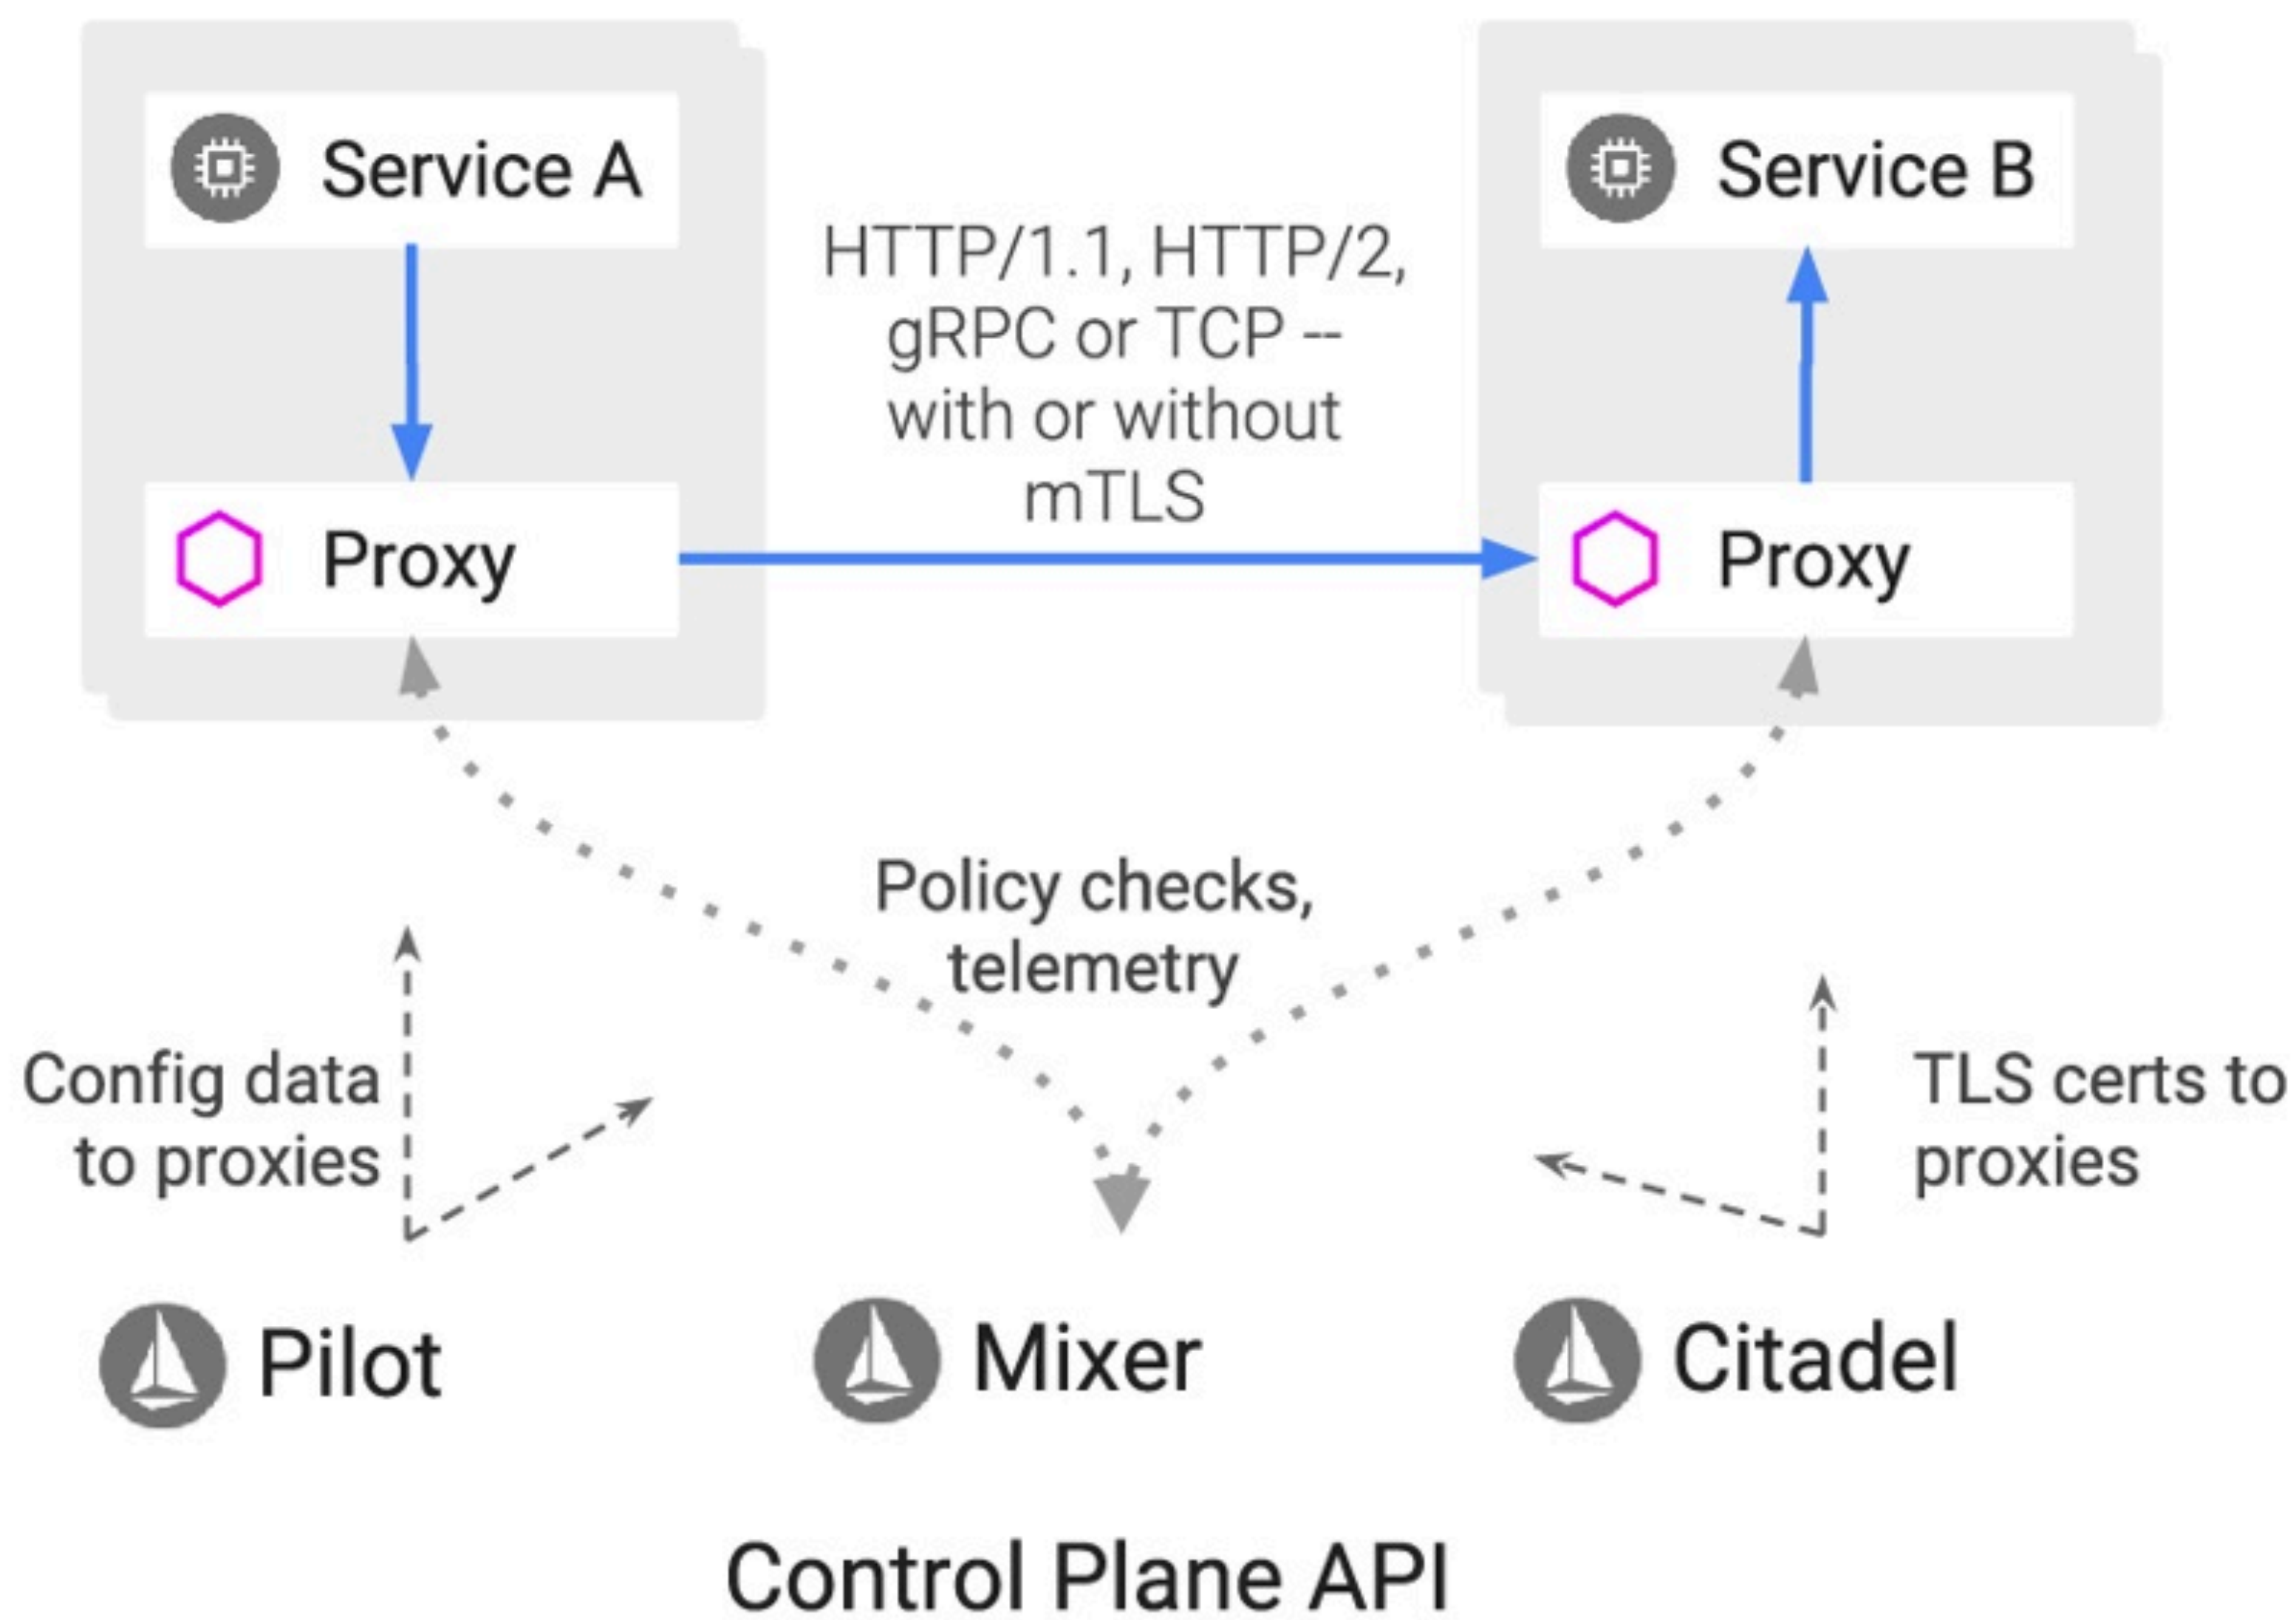
- A service mesh manages all service-to-service communication within a distributed (potentially microservice-based) software system
- Use of “sidecar” proxies that are deployed alongside each service through which all traffic is transparently routed.
- OSI Layer 7 = Communication using HTTP, now any underlying layer like packets TCP, etc
- Dynamic service discovery and traffic management
- Traffic Shadowing for Testing, Traffic Spitting for Canary
- Including but not limited to **Security Enforcement!**
- Linkerd, Istio, Consul, Kuma, Maesh, AWS App Mesh

# OSI Layer Model











# Istio Components

- **Pilot** - Responsible for configuring the Envoy and Mixer at runtime.
- **Proxy / Envoy** - Sidecar proxies per microservice to handle ingress/egress traffic between services in the cluster and from a service to external services.
- **Mixer** - Create a portability layer on top of infrastructure backends. Enforce policies such as ACLs, rate limits, quotas, authentication, request tracing and telemetry collection at an infrastructure level.
- **Citadel / Istio CA** - Secures service to service communication over TLS. Providing a key management system to automate key and certificate generation, distribution, rotation, and revocation.
- **Ingress/Egress** - Configure path based routing for inbound and outbound external traffic.
- **Control Plane API** - Underlying Orchestrator such as Kubernetes or Hashicorp Nomad.

# Mutual TLS Authentication

- TLS Communication and Setup performed through Envoy Proxies
- How Istio handles that traffic:
  - Istio re-routes the outbound traffic from a client to the client's local sidecar Envoy.
  - The client side Envoy starts a mutual TLS handshake with the server side Envoy. During the handshake, the client side Envoy also does a secure naming check to verify that the service account presented in the server certificate is authorized to run the target service.
  - The client side Envoy and the server side Envoy establish a mutual TLS connection, and Istio forwards the traffic from the client side Envoy to the server side Envoy.
  - After authorization, the server side Envoy forwards the traffic to the server service through local TCP connections.



# Thank You



- Email: [dhinojosa@evolutionnext.com](mailto:dhinojosa@evolutionnext.com)
- Github: <https://www.github.com/dhinojosa>
- Twitter: <http://twitter.com/dhinojosa>
- Linked In: <http://www.linkedin.com/in/dhevolutionnext>